

Learned closures and nonlinear manifolds for parametric projection-based reduced models

S. Ares de Parga, Y. Maday

July 15, 2026

Abstract

This work presents a comparative assessment of learned closures and nonlinear manifolds for parametric projection-based reduced-order models of nonlinear time-dependent systems. Within a common reduced-coordinate representation, we compare state-conditioned, parameter–time-conditioned, and hybrid PROM-ANN closures; an intrusive PROM-POD-AE decoder manifold; a non-intrusive POD-NN-ROM direct coefficient map; and a non-intrusive latent POD-DL-ROM. The framework is evaluated on a parametric two-dimensional inviscid Burgers benchmark using both full-residual projection and empirical-cubature hyperreduction.

The study separates three sources of error: the reduced trial representation, the learned-coordinate approximation, and residual hyperreduction. A metric-aware trial basis is designed to limit the transfer of unresolved coordinate errors into residual-solved coordinates. Linear-PROM and linear-HPROM trajectories provide the respective campaign learning targets, which permits data enrichment without acquiring additional high-dimensional snapshots. Controlled dimension and tail-perturbation diagnostics show that a parameter–time closure can feed its secondary-coordinate error back into an intrusive LSPG solve. Enlarging the coefficient training set substantially reduces this effect and improves the intrusive and non-intrusive learned models. The results clarify which error mechanisms are due to the learned manifold and which are introduced by hyperreduction, providing a reproducible basis for selecting closure architectures and training data in projection-based ROMs.

1 Introduction

Projection-based model order reduction (PMOR) is a standard framework for reducing high-dimensional parametric models while preserving governing equations [1, 2]. For nonlinear transport-dominated dynamics, linear subspaces may require large dimensions, which motivates nonlinear manifolds, closure modeling, and latent representations [3, 4, 5].

From a machine-learning viewpoint, four families are relevant to this work:

- intrusive closure-enhanced projection (PROM-ANN) [3],
- intrusive latent-manifold projection (PROM-POD-AE), in the spirit of hyperreduced nonlinear manifold ROMs [6],
- non-intrusive direct-coefficient regression maps $(\boldsymbol{\mu}, t) \mapsto \mathbf{q}_N$,
- non-intrusive latent POD-DL-ROM maps [7].

Naming convention used in this manuscript. When we say **POD-NN-ROM**, we refer to the non-intrusive coefficient-space map $(\boldsymbol{\mu}, t) \mapsto \mathbf{q}_N$. When we say **POD-DL-ROM**, we refer to the non-intrusive latent family in the sense of Fresca–Manzoni [7]. When we say **PROM-POD-AE**, we refer to an intrusive latent-manifold projection route (conceptually aligned with intrusive nonlinear-manifold HPROM literature such as [6]), even though those references use different method names.

The goal is a controlled comparison under strong online compression. All variants use the same HDM residual, time discretization, and fixed trial basis. The PROM experiments evaluate the full residual, whereas the HPROM experiments use model-matched empirical-cubature-method (ECM) rules; this separation exposes the effect of residual sampling instead of conflating it with regression error.

2 Linear PROM/HPROM Reference and Common Basis

2.1 Generic semi-discrete IVP and residual notation

Given the generic, parametric, nonlinear, semi-discrete, first-order initial value problem

$$\mathbf{M}(\boldsymbol{\mu}) \dot{\mathbf{u}}(t; \boldsymbol{\mu}) + \mathbf{f}(\mathbf{u}(t; \boldsymbol{\mu}); \boldsymbol{\mu}) - \mathbf{g}(t; \boldsymbol{\mu}) = \mathbf{0}, \quad \mathbf{u}(0; \boldsymbol{\mu}) = \mathbf{u}_0(\boldsymbol{\mu}), \quad (1)$$

with $\boldsymbol{\mu} \in \mathcal{D} \subset \mathbb{R}^p$, we denote by $\mathbf{r}^{m+1}$ the nonlinear time-discrete residual associated with the selected implicit time integrator. In compact form, each HDM step is written as

$$\mathbf{r}^{m+1}(\mathbf{u}^{m+1}; \boldsymbol{\mu}) = \mathbf{0}. \quad (2)$$

For readability, dependencies on previous time states and Δt are kept implicit in $\mathbf{r}^{m+1}$.

2.2 POD basis from SVD

Let $\mathbf{u}(t; \boldsymbol{\mu}) \in \mathbb{R}^N$ denote the HDM state with parameters $\boldsymbol{\mu} \in \mathcal{D} \subset \mathbb{R}^p$. Define

$$\mathbf{u}_{\text{ref}} := \frac{1}{N_s} \sum_{i=1}^{N_s} \mathbf{u}^i, \quad \mathbf{S}_c = [\mathbf{u}^1 - \mathbf{u}_{\text{ref}} \quad \cdots \quad \mathbf{u}^{N_s} - \mathbf{u}_{\text{ref}}]. \quad (3)$$

Let

$$\mathbf{S}_c = \mathbf{U}_S \boldsymbol{\Sigma}_S \mathbf{Y}_S^T \quad (4)$$

be its thin SVD, and define the linear ROB

$$\mathbf{V}_{\text{tot}} \in \mathbb{R}^{N \times n_{\text{tot}}} \quad (5)$$

as the first n_{tot} left singular vectors of $\mathbf{U}_S$.

The linear reduced reconstruction is

$$\tilde{\mathbf{u}}(t; \boldsymbol{\mu}) = \mathbf{u}_{\text{ref}} + \mathbf{V}_{\text{tot}} \mathbf{q}_N(t; \boldsymbol{\mu}), \quad \mathbf{q}_N(t; \boldsymbol{\mu}) \in \mathbb{R}^{n_{\text{tot}}}. \quad (6)$$

2.3 Linear time-discrete LSPG reference form

Let $\mathbf{r}^{m+1}(\mathbf{u}; \boldsymbol{\mu}) \in \mathbb{R}^N$ denote the HDM time-discrete residual at step $m+1$. The unsampled reference linear LSPG solve is

$$\mathbf{q}_N^{m+1}(\boldsymbol{\mu}) = \arg \min_{\mathbf{x} \in \mathbb{R}^{n_{\text{tot}}}} \|\mathbf{r}^{m+1}(\mathbf{u}_{\text{ref}} + \mathbf{V}_{\text{tot}} \mathbf{x}; \boldsymbol{\mu})\|_2^2. \quad (7)$$

With

$$\mathbf{J}^{m+1,k} := \frac{\partial \mathbf{r}^{m+1}}{\partial \mathbf{u}}, \quad \mathbf{T}_N := \mathbf{V}_{\text{tot}}, \quad \mathbf{W}_N^{m+1,k} := \mathbf{J}^{m+1,k} \mathbf{T}_N, \quad (8)$$

the Gauss–Newton increment solves

$$(\mathbf{W}_N^{m+1,k})^T \mathbf{W}_N^{m+1,k} \Delta \mathbf{q}_N^{m+1,k} = -(\mathbf{W}_N^{m+1,k})^T \mathbf{r}^{m+1}(\mathbf{u}_{\text{ref}} + \mathbf{V}_{\text{tot}} \mathbf{q}_N^{m+1,k}; \boldsymbol{\mu}). \quad (9)$$

2.4 Linear HPROM via the empirical cubature method

Let $\mathcal{E} = \{e_i\}_{i=1}^{N_e}$ denote the set of mesh entities. The time-discrete residual is assembled from local contributions,

$$\mathbf{r}^{m+1}(\mathbf{u}; \boldsymbol{\mu}) = \sum_{e_i \in \mathcal{E}} \mathbf{L}_{e_i}^T \mathbf{r}_{e_i}^{m+1}(\mathbf{L}_{e_i}^+ \mathbf{u}; \boldsymbol{\mu}). \quad (10)$$

Here $\mathbf{r}_{e_i}^{m+1} \in \mathbb{R}^{d_{e_i}}$ is the residual associated with entity e_i , $\mathbf{L}_{e_i} \in \{0, 1\}^{d_{e_i} \times N}$ extracts its residual degrees of freedom, and $\mathbf{L}_{e_i}^+$ extracts the enlarged stencil needed to evaluate that contribution. The latter includes neighboring entities whenever the spatial discretization requires them.

For the affine trial manifold $\boldsymbol{\phi}_N(\mathbf{q}_N) = \mathbf{u}_{\text{ref}} + \mathbf{V}_{\text{tot}} \mathbf{q}_N$, consider Gauss–Newton iteration k at time step $m+1$. With

$$\mathbf{T}_N = \mathbf{V}_{\text{tot}}, \quad \mathbf{J}^{m+1,k} = \left. \frac{\partial \mathbf{r}^{m+1}}{\partial \mathbf{u}} \right|_{\boldsymbol{\phi}_N(\mathbf{q}_N^{m+1,k})}, \quad \mathbf{W}_N^{m+1,k} = \mathbf{J}^{m+1,k} \mathbf{T}_N, \quad (11)$$

the projected LSPG residual has the entity-wise representation

$$\begin{aligned} \mathbf{r}_N^{m+1,k}(\boldsymbol{\mu}) &:= (\mathbf{W}_N^{m+1,k})^T \mathbf{r}^{m+1}(\boldsymbol{\phi}_N(\mathbf{q}_N^{m+1,k}); \boldsymbol{\mu}) \\ &= \sum_{e_i \in \mathcal{E}} \left(\mathbf{L}_{e_i} \mathbf{W}_N^{m+1,k} \right)^T \mathbf{r}_{e_i}^{m+1}(\mathbf{L}_{e_i}^+ \boldsymbol{\phi}_N(\mathbf{q}_N^{m+1,k}); \boldsymbol{\mu}). \end{aligned} \quad (12)$$

Let $\mathbf{J}_{e_i}^{m+1,k} = \partial \mathbf{r}_{e_i}^{m+1} / \partial (\mathbf{L}_{e_i}^+ \mathbf{u})$ at the same state. The Gauss–Newton matrix is assembled analogously,

$$\mathbf{J}_N^{m+1,k} := (\mathbf{W}_N^{m+1,k})^T \mathbf{J}^{m+1,k} \mathbf{T}_N = \sum_{e_i \in \mathcal{E}} \left(\mathbf{L}_{e_i} \mathbf{W}_N^{m+1,k} \right)^T \mathbf{J}_{e_i}^{m+1,k} \left(\mathbf{L}_{e_i}^+ \mathbf{T}_N \right). \quad (13)$$

ECM selects a reduced mesh $\tilde{\mathcal{E}}_N \subset \mathcal{E}$ and positive weights ξ_{N,e_i} that approximate both sums with the same cubature rule:

$$\begin{aligned} \hat{\mathbf{r}}_N^{m+1,k} &= \sum_{e_i \in \tilde{\mathcal{E}}_N} \xi_{N,e_i} \left(\mathbf{L}_{e_i} \mathbf{W}_N^{m+1,k} \right)^T \mathbf{r}_{e_i}^{m+1}(\mathbf{L}_{e_i}^+ \boldsymbol{\phi}_N(\mathbf{q}_N^{m+1,k}); \boldsymbol{\mu}), \\ \hat{\mathbf{J}}_N^{m+1,k} &= \sum_{e_i \in \tilde{\mathcal{E}}_N} \xi_{N,e_i} \left(\mathbf{L}_{e_i} \mathbf{W}_N^{m+1,k} \right)^T \mathbf{J}_{e_i}^{m+1,k} \left(\mathbf{L}_{e_i}^+ \mathbf{T}_N \right). \end{aligned} \quad (14)$$

The sampled Gauss–Newton update is therefore

$$\hat{\mathbf{J}}_N^{m+1,k} \Delta \mathbf{q}_N^{m+1,k} = -\hat{\mathbf{r}}_N^{m+1,k}, \quad \mathbf{q}_N^{m+1,k+1} = \mathbf{q}_N^{m+1,k} + \Delta \mathbf{q}_N^{m+1,k}. \quad (15)$$

Offline cubature construction. At each selected parameter–time state s , the linear reduced coordinate $\mathbf{q}_{N,s}$ and the associated test basis $\mathbf{W}_{N,s}$ define the contribution of entity e_i as

$$\mathbf{c}_{N,e_i}^{(s)} = (\mathbf{L}_{e_i} \mathbf{W}_{N,s})^T \mathbf{r}_{e_i}(\mathbf{L}_{e_i}^+ \boldsymbol{\phi}_N(\mathbf{q}_{N,s}); \boldsymbol{\mu}_s). \quad (16)$$

Stacking these n_{tot} -vectors over the N_{ECM} selected states gives

$$\mathbf{C}_N = \begin{bmatrix} \mathbf{c}_{N,e_1}^{(1)} & \cdots & \mathbf{c}_{N,e_{N_e}}^{(1)} \\ \vdots & & \vdots \\ \mathbf{c}_{N,e_1}^{(N_{\text{ECM}})} & \cdots & \mathbf{c}_{N,e_{N_e}}^{(N_{\text{ECM}})} \end{bmatrix} \in \mathbb{R}^{n_{\text{tot}} N_{\text{ECM}} \times N_e}, \quad \mathbf{b}_N = \mathbf{C}_N \mathbf{1}. \quad (17)$$

The empirical cubature method computes a sparse nonnegative approximation of the full projected residual data,

$$\xi_N = \arg \min_{\xi \geq 0} \|\mathbf{C}_N \xi - \mathbf{b}_N\|_2^2, \quad \tilde{\mathcal{E}}_N = \{e_i \in \mathcal{E} : \xi_{N,e_i} > 0\}. \quad (18)$$

In the present implementation, a direct SVD is used to remove numerically null directions of $\mathbf{C}_N$ before the nonnegative cubature solve. The rule is therefore tied to the reduced manifold and its LSPG test basis. Consequently, the linear rule cannot be reused after replacing the affine trial space by a learned closure or decoder manifold [8, 9, 10, 11].

Shared hyper-reduction principle. The common algebraic step is to assemble entity-wise contributions to the projected LSPG residual and approximate their sum with a sparse positive rule. Several hyperreduction strategies fit this project-then-approximate view. Energy-conserving sampling and weighting (ECSW), for example, acts on the same type of projected contributions [12]. We use ECM throughout this work, following the empirical-cubature construction of [8, 9]; the distinction is methodological, not a change to the reduced manifold or LSPG formulation.

2.5 Basis Metric and Low/High-Coordinate Sensitivity

All learned variants below use a common reduced coordinate system. This subsection records the basis metric used to construct that system and the low/high-coordinate diagnostic that motivates the choice. It is a basis-selection device, not a separate learning model or an additional optimization stage. For the closure constructions introduced later, the common linear basis is partitioned as $\mathbf{V}_{\text{tot}} = [\mathbf{V} \ \bar{\mathbf{V}}]$, with online coordinates $\mathbf{q}$ and reconstructed coordinates $\bar{\mathbf{q}}$. This partition is not only an approximation choice: in an intrusive LSPG solve, errors in the reconstructed high-order coordinates can feed back into the low-order coordinates through the residual normal equations. The same $\mathbf{V}_{\text{tot}}$ also defines the coordinate system used by the decoder and direct regression models. Consequently, the metric used to construct the basis is part of the common experimental setup, rather than a model-specific preprocessing choice.

2.5.1 Euclidean and M-regularized LSPG-sensitive metrics

For comparison, the Euclidean POD basis is obtained from the centered snapshot matrix $\mathbf{S}_c$. We also consider a residual-sensitive metric that combines a physical mass contribution with an averaged LSPG sensitivity:

$$\mathbf{W} = \varepsilon \mathbf{M} + \sum_{s \in \mathcal{S}_{\text{met}}} \omega_s (\mathbf{J}_s)^T \mathbf{P}_s \mathbf{J}_s. \quad (19)$$

Here $\mathbf{M}$ is the block mass/area matrix associated with the state inner product, $\mathbf{J}_s = \partial \mathbf{r}_s / \partial \mathbf{u}$ is the time-discrete residual Jacobian at a selected snapshot, $\mathbf{P}_s$ is the nonnegative sampling/weighting operator used for the residual metric, and ω_s are quadrature and parameter weights. The small positive parameter ε regularizes the metric and preserves physical state weighting when the residual sensitivity is weak or locally rank deficient.

The corresponding weighted POD basis is obtained from the generalized correlation matrix

$$\mathbf{C} = \mathbf{S}_c^T \mathbf{W} \mathbf{S}_c. \quad (20)$$

If $\mathbf{C} \mathbf{y}_i = \lambda_i \mathbf{y}_i$, the associated modes are

$$\phi_i = \frac{1}{\sqrt{\lambda_i}} \mathbf{S}_c \mathbf{y}_i, \quad \phi_i^T \mathbf{W} \phi_j = \delta_{ij}. \quad (21)$$

In practice, the metric is assembled on a representative subset of parameter and time samples, selected independently from the final test points. The resulting basis is then used uniformly by all reduced models; it is not tuned separately for a particular closure map. The Burgers-specific sampling and regularization values are given with the experimental protocol in Section 7.

2.5.2 Low/high transfer diagnostic

To quantify the potential effect of prescribing a secondary-coordinate block, consider a local linearization of the LSPG normal equations around a reference reduced trajectory. Let

$$\mathbf{A}_s = (\mathbf{J}_s)^T \mathbf{P}_s \mathbf{J}_s, \quad \mathbf{H}_s = \begin{bmatrix} \mathbf{V}^T \mathbf{A}_s \mathbf{V} & \mathbf{V}^T \mathbf{A}_s \bar{\mathbf{V}} \\ \bar{\mathbf{V}}^T \mathbf{A}_s \mathbf{V} & \bar{\mathbf{V}}^T \mathbf{A}_s \bar{\mathbf{V}} \end{bmatrix}. \quad (22)$$

If the high coordinates are perturbed by $\delta \bar{\mathbf{q}}$ and $\mathbf{H}_{LL}$ is nonsingular, the corresponding first-order low-coordinate response is approximated by

$$\delta \mathbf{q} \approx -(\mathbf{H}_{LL})^{-1} \mathbf{H}_{LH} \delta \bar{\mathbf{q}} = \mathbf{T}_{LH} \delta \bar{\mathbf{q}}. \quad (23)$$

Thus, $\mathbf{T}_{LH}$ identifies high-coordinate directions that are most likely to perturb the resolved LSPG coordinates. The purpose of the M -regularized LSPG-sensitive metric is not to optimize one regression map, but to construct a basis for which this low/high transfer is less amplifying without compromising the linear reduced representation. The closure maps that use this partition are defined formally in Section 3.

2.5.3 Basis-selection criterion

The metric is selected at the level of the common reduced coordinate system, before choosing a closure architecture or a non-intrusive regressor. It therefore changes not only the trial space $\mathbf{V}_{\text{tot}}$, but also the reduced-coordinate trajectories used as learning targets. Selecting the metric by the validation loss of one learned map would consequently confound basis quality with the capacity and optimization of that particular map.

We instead impose two basis-level requirements. First, at the retained dimension n_{tot} , the metric must preserve the state accuracy of the full-coordinate linear LSPG reference. This prevents a reduction in low/high transfer from being obtained simply by degrading the linear trial representation. The ECM rules used subsequently are constructed only after this basis has been fixed and are matched to the relevant trial manifold. Second, for a partition $\mathbf{q}_N = [\mathbf{q}; \bar{\mathbf{q}}]$, it should reduce the amplification measured by (23). This second requirement is relevant whenever a reduced model supplies part of the coordinate vector rather than solving all n_{tot} coordinates through the residual; the specific models considered here are defined in the following sections.

The direct transfer diagnostic and the complementary oracle perturbation test are reported in Appendix B. Together, they show that the M -regularized LSPG-sensitive basis reduces the high-to-low coupling while leaving the linear reference essentially unchanged. Accordingly, unless stated otherwise, every model below uses this single coordinate basis; the intrusive variants use it as their trial basis. This common choice is essential: differences among the subsequent models can then be attributed to the learned coordinate relation and, for HPROMs, to the model-matched cubature rule rather than to a model-specific trial space.

3 Method I: PROM-ANN Closure Models

3.1 Manifold definitions

The PROM-ANN construction considered here is based on neural-network-augmented projection [3], in which a nonlinear trial manifold augments a retained linear subspace with a learned

closure for the remaining reduced coordinates. In its original state-conditioned form, the learned relation supplies discarded reduced coordinates while an intrusive projection solve retains responsibility for the primary block [3]. This separation is not tied to a neural-network regressor: Gaussian-process and radial-basis-function closures have been considered for the same latent closure role [11]. Related assessments of projection-based reduction for hypersonic-flow benchmarks provide an application context in which retaining the projection solve is particularly relevant [13, 14].

Let the common linear trial basis and its coordinates be partitioned as

$$\mathbf{V}_{\text{tot}} = [\mathbf{V} \ \bar{\mathbf{V}}], \quad \mathbf{q}_N = \begin{bmatrix} \mathbf{q} \\ \bar{\mathbf{q}} \end{bmatrix}, \quad n + \bar{n} = n_{\text{tot}}. \quad (24)$$

The associated linear reduced state is therefore

$$\mathbf{u}(t; \boldsymbol{\mu}) \approx \mathbf{u}_{\text{ref}} + \mathbf{V} \mathbf{q}(t; \boldsymbol{\mu}) + \bar{\mathbf{V}} \bar{\mathbf{q}}(t; \boldsymbol{\mu}). \quad (25)$$

The first block $\mathbf{q} \in \mathbb{R}^n$ contains the coordinates retained as online unknowns; the secondary block $\bar{\mathbf{q}} \in \mathbb{R}^{\bar{n}}$ contains the remaining coordinates.

The free/dependent coordinate split in (24) also appears in nonlinear compressive reduced-basis theory. For a solution manifold parameterized by p physical parameters, the local sensing analysis identifies $n = p$ as an appropriate local number of free coordinates and supports a small excess above p over a wider parameter region [5]. This result motivates treating n as a quantity to be assessed rather than fixing it only from POD energy. It is not, however, a direct prescription for the transient LSPG setting below, where the learned relation can also depend on time and the online dynamics. In particular, the neural and GPR/RBF closure constructions above do not select n from the physical parameter dimension; they prescribe how the dependent coordinates are represented once a split has been chosen [3, 11].

PROM-ANN replaces the secondary block by a learned map

$$\bar{\mathbf{q}} = \mathcal{F}(\mathbf{q}, \boldsymbol{\mu}, t), \quad \mathcal{F} : \mathbb{R}^n \times \mathcal{D} \times [0, T] \longrightarrow \mathbb{R}^{\bar{n}}, \quad (26)$$

which defines the parameter-time-dependent nonlinear trial manifold

$$\mathcal{M}_{\mathcal{F}}(\boldsymbol{\mu}, t) = \{ \mathbf{u}_{\text{ref}} + \mathbf{V} \mathbf{q} + \bar{\mathbf{V}} \mathcal{F}(\mathbf{q}, \boldsymbol{\mu}, t) : \mathbf{q} \in \mathbb{R}^n \}. \quad (27)$$

Thus, the network does not approximate the high-dimensional state directly: it supplies only the secondary reduced coordinates, while the primary coordinates are determined by the residual minimization described below.

The following three choices are information patterns within this common closure parameterization, not three unrelated nonlinear-manifold methods. Case 1 is the state-conditioned PROM-ANN closure described above. Cases 2 and 3 are controlled variants used here to separate a purely parameter-time-prescribed tail from a tail that can also respond to the current primary iterate. The three maps are

$$\text{Case 1: } \bar{\mathbf{q}} = \mathcal{N}(\mathbf{q}), \quad (28)$$

$$\text{Case 2: } \bar{\mathbf{q}} = \mathcal{M}(\boldsymbol{\mu}, t), \quad (29)$$

$$\text{Case 3: } \bar{\mathbf{q}} = \mathcal{H}(\mathbf{q}, \boldsymbol{\mu}, t). \quad (30)$$

Case 1 is a state-conditioned closure: the secondary coordinates are inferred only from the currently solved primary coordinates. Case 2 is a parameter-time-conditioned closure: at every time step it injects a secondary-coordinate vector predicted from $(\boldsymbol{\mu}, t)$, independently of the current primary iterate. Case 3 combines both sources of information. Consequently, Case 2 has the simplest online tangent, whereas Cases 1 and 3 can adapt the closure as the primary coordinates change during the nonlinear solve.

In the n_{tot} -dimensional linear PROM or linear HPROM reference used for target generation, both $\mathbf{q}$ and $\bar{\mathbf{q}}$ are solved online. In the PROM-ANN models, only $\mathbf{q}$ remains an online unknown; the secondary coordinates are evaluated from the learned closure. This is an n -dimensional nonlinear manifold embedded in the same n_{tot} -dimensional coefficient space used to generate the training targets.

The three cases differ in how the closure interacts with the LSPG tangent. In Case 1 and Case 3 the tangent contains $\bar{\mathbf{V}} \partial \mathcal{F} / \partial \mathbf{q}$, whereas in Case 2 the closure is parameter-time conditioned and the tangent with respect to the online coordinate is simply $\mathbf{V}$. Case 2 is therefore the cheapest intrusive closure, but it is also the most exposed to high-coordinate injection errors because the injected block is not corrected through a state-dependent tangent.

3.2 LSPG/Gauss-Newton solve

At each time step,

$$\mathbf{q}^{m+1}(\boldsymbol{\mu}) = \arg \min_{\mathbf{x} \in \mathbb{R}^n} \|\mathbf{r}^{m+1}(\tilde{\mathbf{u}}^{m+1}(\mathbf{x}; \boldsymbol{\mu}); \boldsymbol{\mu})\|_2^2. \quad (31)$$

Define

$$\mathbf{J}^{m+1,k} := \frac{\partial \mathbf{r}^{m+1}}{\partial \mathbf{u}}, \quad \mathbf{T}_q^{m+1,k} := \frac{\partial \tilde{\mathbf{u}}^{m+1}}{\partial \mathbf{q}} = \mathbf{V} + \bar{\mathbf{V}} \frac{\partial \mathcal{F}}{\partial \mathbf{q}}, \quad \mathbf{W}_q^{m+1,k} = \mathbf{J}^{m+1,k} \mathbf{T}_q^{m+1,k}. \quad (32)$$

The Gauss-Newton increment solves

$$(\mathbf{W}_q^{m+1,k})^T \mathbf{W}_q^{m+1,k} \Delta \mathbf{q}^{m+1,k} = -(\mathbf{W}_q^{m+1,k})^T \mathbf{r}^{m+1}(\tilde{\mathbf{u}}^{m+1}(\mathbf{q}^{m+1,k}; \boldsymbol{\mu}); \boldsymbol{\mu}). \quad (33)$$

3.3 ECM hyperreduced form

For a closure choice $c \in \{1, 2, 3\}$, define the learned trial manifold and its tangent by

$$\phi_c(\mathbf{q}; \boldsymbol{\mu}, t) = \mathbf{u}_{\text{ref}} + \mathbf{V} \mathbf{q} + \bar{\mathbf{V}} \mathcal{F}_c(\mathbf{q}, \boldsymbol{\mu}, t), \quad \mathbf{T}_c = \frac{\partial \phi_c}{\partial \mathbf{q}}, \quad \mathbf{W}_c = \mathbf{J} \mathbf{T}_c. \quad (34)$$

Here $\mathcal{F}_c$ is respectively $\mathcal{N}$, $\mathcal{M}$, or $\mathcal{H}$. In particular, $\mathbf{T}_2 = \mathbf{V}$, whereas $\mathbf{T}_1$ and $\mathbf{T}_3$ include the neural-network derivative. This manifold-dependent tangent is part of the LSPG test basis and must therefore also be included in the ECM construction.

At a given Gauss-Newton iterate $(\mathbf{q}^{m+1,k}, \boldsymbol{\mu}, t^{m+1})$, the projected LSPG residual has the same entity-wise form as in Eq. (12),

$$\begin{aligned} \mathbf{r}_c^{m+1}(\mathbf{q}; \boldsymbol{\mu}) &:= \mathbf{W}_c^T \mathbf{r}^{m+1}(\phi_c(\mathbf{q}; \boldsymbol{\mu}, t^{m+1}); \boldsymbol{\mu}) \\ &= \sum_{e_i \in \mathcal{E}} (\mathbf{L}_{e_i} \mathbf{W}_c)^T \mathbf{r}_{e_i}^{m+1}(\mathbf{L}_{e_i}^+ \phi_c(\mathbf{q}; \boldsymbol{\mu}, t^{m+1}); \boldsymbol{\mu}). \end{aligned} \quad (35)$$

Writing $\mathbf{J}_{e_i} = \partial \mathbf{r}_{e_i} / \partial (\mathbf{L}_{e_i}^+ \mathbf{u})$ at the current state, the corresponding Gauss-Newton matrix is

$$\mathbf{J}_c^{m+1}(\mathbf{q}; \boldsymbol{\mu}) = \sum_{e_i \in \mathcal{E}} (\mathbf{L}_{e_i} \mathbf{W}_c)^T \mathbf{J}_{e_i} (\mathbf{L}_{e_i}^+ \mathbf{T}_c). \quad (36)$$

Thus, ECM represents the entity contributions to the LSPG normal equations; it does not approximate the neural network itself.

Offline construction of the cubature rule. For every selected parameter–time training state, the HDM state $\mathbf{u}_s^{m_s+1}$ is first represented on the trained manifold. The fitted coordinate is

$$\mathbf{q}_{c,s} = \arg \min_{\mathbf{y} \in \mathbb{R}^n} \|\mathbf{u}_s^{m_s+1} - \phi_c(\mathbf{y}; \boldsymbol{\mu}_s, t_s^{m_s+1})\|_2^2. \quad (37)$$

For Case 2, this is the linear least-squares problem obtained by subtracting the prescribed tail,

$$\mathbf{q}_{2,s} = \mathbf{V}^\dagger [\mathbf{u}_s^{m_s+1} - \mathbf{u}_{\text{ref}} - \bar{\mathbf{V}}\mathcal{M}(\boldsymbol{\mu}_s, t_s^{m_s+1})], \quad (38)$$

where $\dagger$ denotes the Moore–Penrose inverse. For Cases 1 and 3, define the state mismatch and its tangent by

$$\boldsymbol{\delta}_{c,s}(\mathbf{y}) = \phi_c(\mathbf{y}; \boldsymbol{\mu}_s, t_s^{m_s+1}) - \mathbf{u}_s^{m_s+1}, \quad \mathbf{T}_{c,s}(\mathbf{y}) = \frac{\partial \phi_c}{\partial \mathbf{y}}(\mathbf{y}; \boldsymbol{\mu}_s, t_s^{m_s+1}), \quad c \in \{1, 3\}. \quad (39)$$

With $\mathbf{P}_{\leq n} = [\mathbf{I}_n \mathbf{0}]$, we initialize from the primary block of the full linear least-squares coordinates and apply Gauss–Newton updates,

$$\mathbf{q}_{c,s}^{(0)} = \mathbf{P}_{\leq n} \mathbf{V}_{\text{tot}}^\dagger (\mathbf{u}_s^{m_s+1} - \mathbf{u}_{\text{ref}}), \quad \mathbf{q}_{c,s}^{(k+1)} = \mathbf{q}_{c,s}^{(k)} - \mathbf{T}_{c,s}(\mathbf{q}_{c,s}^{(k)})^\dagger \boldsymbol{\delta}_{c,s}(\mathbf{q}_{c,s}^{(k)}). \quad (40)$$

The converged iterate defines $\mathbf{q}_{c,s}$ in (37); in practice the update is solved as a least-squares Gauss–Newton step. Evaluating (35) at this fitted state gives the entity contribution

$$\mathbf{c}_{c,e_i}^{(s)} = (\mathbf{L}_{e_i} \mathbf{W}_{c,s})^T \mathbf{r}_{e_i}(\mathbf{L}_{e_i}^+ \phi_c(\mathbf{q}_{c,s}; \boldsymbol{\mu}_s, t_s); \boldsymbol{\mu}_s) \in \mathbb{R}^n. \quad (41)$$

Stacking these vectors over the N_{ECM} training states gives the ECM contribution matrix

$$\mathbf{C}_c = \begin{bmatrix} \mathbf{c}_{c,e_1}^{(1)} & \cdots & \mathbf{c}_{c,e_{N_e}}^{(1)} \\ \vdots & & \vdots \\ \mathbf{c}_{c,e_1}^{(N_{\text{ECM}})} & \cdots & \mathbf{c}_{c,e_{N_e}}^{(N_{\text{ECM}})} \end{bmatrix} \in \mathbb{R}^{n N_{\text{ECM}} \times N_e}, \quad \mathbf{b}_c = \mathbf{C}_c \mathbf{1}. \quad (42)$$

Each column of $\mathbf{C}_c$ is consequently the complete contribution of one mesh entity to all selected projected residuals. ECM seeks a sparse positive cubature rule by solving

$$\boldsymbol{\xi}_c = \arg \min_{\boldsymbol{\xi} \geq \mathbf{0}} \|\mathbf{C}_c \boldsymbol{\xi} - \mathbf{b}_c\|_2^2, \quad \tilde{\mathcal{E}}_c = \{e_i \in \mathcal{E} : \xi_{c,e_i} > 0\}. \quad (43)$$

In the present implementation, a deterministic direct SVD of $\mathbf{C}_c$ first identifies its numerical range, and ECM determines the nonnegative weights in that compressed space. The parameter–time sample set, the trained closure, and its tangent are therefore all part of the rule; a linear-HPROM rule or another closure’s rule is not reused [8, 9, 10, 11].

Online HPROM solve. Only the entities in $\tilde{\mathcal{E}}_c$ and the neighboring entities required by their discrete stencils are evaluated online. The sampled LSPG normal equations are

$$\begin{aligned} \hat{\mathbf{r}}_c^{m+1} &= \sum_{e_i \in \tilde{\mathcal{E}}_c} \xi_{c,e_i} (\mathbf{L}_{e_i} \mathbf{W}_c)^T \mathbf{r}_{e_i}^{m+1}(\mathbf{L}_{e_i}^+ \phi_c(\mathbf{q}; \boldsymbol{\mu}, t^{m+1}); \boldsymbol{\mu}), \\ \hat{\mathbf{J}}_c^{m+1} &= \sum_{e_i \in \tilde{\mathcal{E}}_c} \xi_{c,e_i} (\mathbf{L}_{e_i} \mathbf{W}_c)^T \mathbf{J}_{e_i}(\mathbf{L}_{e_i}^+ \mathbf{T}_c). \end{aligned} \quad (44)$$

The nonlinear HPROM solves $\hat{\mathbf{r}}_c^{m+1} = \mathbf{0}$ by the sampled Gauss–Newton update

$$\hat{\mathbf{J}}_c^{m+1,k} \Delta \mathbf{q}^{m+1,k} = -\hat{\mathbf{r}}_c^{m+1,k}, \quad \mathbf{q}^{m+1,k+1} = \mathbf{q}^{m+1,k} + \Delta \mathbf{q}^{m+1,k}. \quad (45)$$

Thus ECM reduces the residual and Jacobian evaluations, while automatic differentiation of the closure continues to provide the nonlinear-manifold tangent in the sampled Gauss–Newton solve.

Offline training targets. For Cases 1–3, supervised data are extracted from the corresponding linear solver trajectory $\mathbf{q}_N = [\mathbf{q}; \bar{\mathbf{q}}]$: the linear PROM for the PROM campaign and the linear HPROM for the HPROM campaign. Case 1 trains $\mathcal{N}$ from $(\mathbf{q}, \bar{\mathbf{q}})$ and Case 3 trains $\mathcal{H}$ from $(\mathbf{q}, \boldsymbol{\mu}, t, \bar{\mathbf{q}})$. The practical Case 2 training choice used in the reported comparisons is specified in the training protocol in Sec. 7.3.

4 Method II: PROM-POD-AE Latent Manifold

4.1 Autoencoder and latent manifold

Unlike Method I, which learns only the secondary coordinate block, the autoencoder acts on the complete linear reduced-coordinate vector $\mathbf{q}_N$. It replaces that vector by a low-dimensional latent coordinate $\mathbf{z}$, which is the sole online unknown of the intrusive solve. Let

$$\mathbf{z} = \mathcal{E}(\mathbf{q}_N), \quad \hat{\mathbf{q}}_N = \mathcal{D}(\mathbf{z}), \quad \mathbf{z} \in \mathbb{R}^{n_z}, \quad n_z \ll n_{\text{tot}}. \quad (46)$$

The intrusive latent manifold is

$$\tilde{\mathbf{u}}(t; \boldsymbol{\mu}) = \mathbf{u}_{\text{ref}} + \mathbf{V}_{\text{tot}} \mathcal{D}(\mathbf{z}(t; \boldsymbol{\mu})). \quad (47)$$

This is an intrusive latent-manifold projection approach. It is aligned with intrusive nonlinear-manifold HPROM literature [6], but is distinct from the non-intrusive POD-DL-ROM line [7].

Offline autoencoder training. Given campaign-consistent linear-solver targets $\mathbf{q}_N^{\text{ref}}$, the autoencoder $(\mathcal{E}, \mathcal{D})$ is trained by reconstruction in reduced coordinates:

$$\mathcal{L}_{\text{AE}} = \frac{1}{N_{\text{samp}}} \sum_{j=1}^{N_{\text{samp}}} \left\| \mathbf{q}_{N,j}^{\text{ref}} - \mathcal{D}(\mathcal{E}(\mathbf{q}_{N,j}^{\text{ref}})) \right\|_2^2. \quad (48)$$

The resulting decoder $\mathcal{D}$ defines the nonlinear trial manifold used in the intrusive online LSPG solve (49).

4.2 LSPG/Gauss–Newton solve in latent coordinates

The online unknown is $\mathbf{z}^{m+1} \in \mathbb{R}^{n_z}$:

$$\mathbf{z}^{m+1}(\boldsymbol{\mu}) = \arg \min_{\mathbf{y} \in \mathbb{R}^{n_z}} \left\| \mathbf{r}^{m+1}(\mathbf{u}_{\text{ref}} + \mathbf{V}_{\text{tot}} \mathcal{D}(\mathbf{y}); \boldsymbol{\mu}) \right\|_2^2. \quad (49)$$

Define

$$\mathbf{T}_z^{m+1,k} := \frac{\partial \tilde{\mathbf{u}}^{m+1}}{\partial \mathbf{z}} = \mathbf{V}_{\text{tot}} \frac{\partial \mathcal{D}}{\partial \mathbf{z}}, \quad \mathbf{W}_z^{m+1,k} = \mathbf{J}^{m+1,k} \mathbf{T}_z^{m+1,k}. \quad (50)$$

The Gauss–Newton step solves

$$(\mathbf{W}_z^{m+1,k})^T \mathbf{W}_z^{m+1,k} \Delta \mathbf{z}^{m+1,k} = -(\mathbf{W}_z^{m+1,k})^T \mathbf{r}^{m+1}(\tilde{\mathbf{u}}^{m+1}(\mathbf{z}^{m+1,k}; \boldsymbol{\mu}); \boldsymbol{\mu}). \quad (51)$$

4.3 ECM hyperreduced form

The POD-AE HPROM uses the same sampled-LSPG/ECM construction, but its nonlinear coordinate map is a decoder of the complete reduced state rather than a closure for a secondary block. Its manifold, tangent, and LSPG test basis are

$$\phi_{\text{AE}}(\mathbf{z}) = \mathbf{u}_{\text{ref}} + \mathbf{V}_{\text{tot}} \mathcal{D}(\mathbf{z}), \quad \mathbf{T}_{\text{AE}} = \frac{\partial \phi_{\text{AE}}}{\partial \mathbf{z}} = \mathbf{V}_{\text{tot}} \frac{\partial \mathcal{D}}{\partial \mathbf{z}}, \quad \mathbf{W}_{\text{AE}} = \mathbf{J} \mathbf{T}_{\text{AE}}. \quad (52)$$

The decoder derivative varies with the latent state. Consequently, neither the affine linear-HPROM test basis nor a PROM–ANN rule represents the POD-AE manifold.

Offline construction for the decoder manifold. For each ECM training state, the latent coordinate is obtained by fitting the state on the trained decoder manifold,

$$\mathbf{z}_s = \arg \min_{\mathbf{y} \in \mathbb{R}^{n_z}} \|\mathbf{u}_s^{m_s+1} - \phi_{\text{AE}}(\mathbf{y})\|_2^2, \quad (53)$$

which is also a nonlinear least-squares problem. With $\mathbf{q}_{N,s}^{(0)} = \mathbf{V}_{\text{tot}}^\dagger (\mathbf{u}_s^{m_s+1} - \mathbf{u}_{\text{ref}})$, the encoder supplies the initial latent coordinate and Gauss–Newton refines it as

$$\mathbf{z}_s^{(0)} = \mathcal{E}(\mathbf{q}_{N,s}^{(0)}), \quad \mathbf{z}_s^{(k+1)} = \mathbf{z}_s^{(k)} - \mathbf{T}_{\text{AE}}(\mathbf{z}_s^{(k)})^\dagger \left[\phi_{\text{AE}}(\mathbf{z}_s^{(k)}) - \mathbf{u}_s^{m_s+1} \right]. \quad (54)$$

The converged iterate is the $\mathbf{z}_s$ used below. At this fitted latent state, the POD-AE entity contribution is

$$\mathbf{c}_{\text{AE},e_i}^{(s)} = (\mathbf{L}_{e_i} \mathbf{W}_{\text{AE},s})^T \mathbf{r}_{e_i}(\mathbf{L}_{e_i}^\dagger \phi_{\text{AE}}(\mathbf{z}_s); \boldsymbol{\mu}_s). \quad (55)$$

The matrix $\mathbf{C}_{\text{AE}}$ is formed by stacking $\mathbf{c}_{\text{AE},e_i}^{(s)}$ exactly as in Eq. (42), and $\mathbf{b}_{\text{AE}} = \mathbf{C}_{\text{AE}} \mathbf{1}$. The weights $\boldsymbol{\xi}_{\text{AE}}$ solve the same nonnegative cubature problem in Eq. (43). A separate direct SVD and ECM calculation are therefore required for the POD-AE decoder: its projected residual contributions depend on the decoder Jacobian at the manifold-fitted training states.

Online POD-AE HPROM solve. Let $\tilde{\mathcal{E}}_{\text{AE}}$ be the support of $\boldsymbol{\xi}_{\text{AE}}$. The sampled latent LSPG normal equations are

$$\begin{aligned} \hat{\mathbf{r}}_{\text{AE}}^{m+1} &= \sum_{e_i \in \tilde{\mathcal{E}}_{\text{AE}}} \xi_{\text{AE},e_i} (\mathbf{L}_{e_i} \mathbf{W}_{\text{AE}})^T \mathbf{r}_{e_i}^{m+1}(\mathbf{L}_{e_i}^\dagger \phi_{\text{AE}}(\mathbf{z}); \boldsymbol{\mu}), \\ \hat{\mathbf{J}}_{\text{AE}}^{m+1} &= \sum_{e_i \in \tilde{\mathcal{E}}_{\text{AE}}} \xi_{\text{AE},e_i} (\mathbf{L}_{e_i} \mathbf{W}_{\text{AE}})^T \mathbf{J}_{e_i}(\mathbf{L}_{e_i}^\dagger \mathbf{T}_{\text{AE}}). \end{aligned} \quad (56)$$

The latent coordinates are advanced by the sampled Gauss–Newton update,

$$\hat{\mathbf{J}}_{\text{AE}}^{m+1,k} \Delta \mathbf{z}^{m+1,k} = -\hat{\mathbf{r}}_{\text{AE}}^{m+1,k}, \quad \mathbf{z}^{m+1,k+1} = \mathbf{z}^{m+1,k} + \Delta \mathbf{z}^{m+1,k}, \quad (57)$$

Thus the online complexity is reduced through local residual and Jacobian evaluation, whereas the latent manifold itself and its tangent are retained exactly through the trained decoder and automatic differentiation.

Comparison of intrusive trial maps. Table 1 collects the trial decoder and tangent basis for every intrusive model considered above. It makes explicit that the models differ in their online coordinates and trial maps, not in the governing time-discrete residual.

Table 1: Trial decoders and tangent bases of the intrusive models. The linear PROM/HPROM row uses the complete linear coordinate vector. The remaining rows solve only their displayed online variable.

Model	Online variable	Trial decoder ϕ	Tangent $\mathbf{T}$
Linear PROM/HPROM	$\mathbf{q}_N$	$\mathbf{u}_{\text{ref}} + \mathbf{V}_{\text{tot}} \mathbf{q}_N$	$\mathbf{V}_{\text{tot}}$
PROM-ANN Case 1	$\mathbf{q}$	$\mathbf{u}_{\text{ref}} + \mathbf{V} \mathbf{q} + \bar{\mathbf{V}} \mathcal{N}(\mathbf{q})$	$\mathbf{V} + \bar{\mathbf{V}} \frac{\partial \mathcal{N}}{\partial \mathbf{q}}$
PROM-ANN Case 2	$\mathbf{q}$	$\mathbf{u}_{\text{ref}} + \mathbf{V} \mathbf{q} + \bar{\mathbf{V}} \mathcal{M}(\boldsymbol{\mu}, t)$	$\mathbf{V}$
PROM-ANN Case 3	$\mathbf{q}$	$\mathbf{u}_{\text{ref}} + \mathbf{V} \mathbf{q} + \bar{\mathbf{V}} \mathcal{H}(\mathbf{q}, \boldsymbol{\mu}, t)$	$\mathbf{V} + \bar{\mathbf{V}} \frac{\partial \mathcal{H}}{\partial \mathbf{q}}$
PROM-POD-AE	$\mathbf{z}$	$\mathbf{u}_{\text{ref}} + \mathbf{V}_{\text{tot}} \mathcal{D}(\mathbf{z})$	$\mathbf{V}_{\text{tot}} \frac{\partial \mathcal{D}}{\partial \mathbf{z}}$

Common online principle. For every row in Table 1, the online state is obtained by minimizing the same time-discrete residual over the displayed trial map. In a PROM this residual is evaluated in full; in an HPROM its corresponding projected normal equations are approximated by the model-specific ECM rule. The non-intrusive POD–NN-ROM and POD–DL-ROM maps are excluded because they perform no online residual minimization.

5 Method III: Non-Intrusive POD–NN-ROM

Model definition. The non-intrusive baseline maps exogenous inputs directly to reduced coefficients:

$$\mathbf{q}_N(t; \boldsymbol{\mu}) \approx \mathcal{G}_q(\boldsymbol{\mu}, t). \quad (58)$$

It therefore predicts both primary and secondary coordinates; no subset of $\mathbf{q}_N$ is corrected by an online residual minimization.

Training objective and online cost. Given the campaign-consistent linear-solver targets $\mathbf{q}_N^{\text{ref}}(t; \boldsymbol{\mu})$, one minimizes a supervised loss (e.g., MSE)

$$\mathcal{L}_{\text{data}} = \frac{1}{N_{\text{samp}}} \sum_{j=1}^{N_{\text{samp}}} \left\| \mathcal{G}_q(\boldsymbol{\mu}_j, t_j) - \mathbf{q}_{N,j}^{\text{ref}} \right\|_2^2. \quad (59)$$

For the PROM campaign, $\mathbf{q}_N^{\text{ref}}$ is the linear-PROM LSPG trajectory; for the HPROM campaign, it is the corresponding linear-HPROM trajectory. This differs intentionally from the usual non-intrusive practice of using projected HDM snapshots as labels. Online POD–NN-ROM prediction is a direct forward evaluation of $\mathcal{G}_q$ followed by reconstruction with $\mathbf{V}_{\text{tot}}$.

6 Method IV: POD–DL-ROM (Latent, Non-Intrusive)

Model definition. The second non-intrusive baseline follows the POD–DL-ROM family [7]:

$$\mathbf{z}(t; \boldsymbol{\mu}) \approx \mathcal{G}_z(\boldsymbol{\mu}, t), \quad \mathbf{q}_N(t; \boldsymbol{\mu}) \approx \mathcal{D}(\mathcal{G}_z(\boldsymbol{\mu}, t)), \quad (60)$$

where $\mathcal{D}$ is a decoder trained offline and $\mathbf{z} \in \mathbb{R}^{n_z}$. As in Method III, this is a direct exogenous-input map: neither $\mathbf{z}$ nor any component of $\mathbf{q}_N$ is solved by LSPG online.

Training objective and online cost. In the present formulation, the autoencoder block $(\mathcal{E}, \mathcal{D})$ follows the same training strategy and normalization used in Section 4: reconstruction of the corresponding linear-solver $\mathbf{q}_N$ snapshots through a latent bottleneck. The POD–DL-ROM dynamics map then replaces the POD–NN-ROM output variable $\mathbf{q}_N$ by $\mathbf{z}$, i.e., it learns $(\boldsymbol{\mu}, t) \mapsto \mathbf{z}$ and composes with the same decoder $\mathcal{D}$:

$$\mathbf{q}_N(\boldsymbol{\mu}, t) \approx \mathcal{D}(\mathcal{G}_z(\boldsymbol{\mu}, t)). \quad (61)$$

Hence, POD–NN-ROM and POD–DL-ROM share the same exogenous-input principle $(\boldsymbol{\mu}, t) \mapsto (\cdot)$, with POD–DL-ROM acting in latent coordinates as proposed in [7]. Online evaluation remains non-intrusive (single forward pass) and does not involve residual minimization. Its autoencoder and latent-dynamics targets follow the same linear-PROM or linear-HPROM teacher convention as POD–NN-ROM; they are not projections of HDM snapshots.

7 Burgers Benchmark and Experimental Protocol

7.1 Governing equations and parameter domain

We use the 2D parametric inviscid Burgers benchmark from [3], in conservative vector form:

$$\frac{\partial \mathbf{U}}{\partial t} + \frac{\partial \mathbf{F}(\mathbf{U})}{\partial x} + \frac{\partial \mathbf{G}(\mathbf{U})}{\partial y} = \mathbf{S}(x; \boldsymbol{\mu}), \quad \mathbf{U} = \begin{bmatrix} u_x \\ u_y \end{bmatrix}, \quad (62)$$

with

$$\mathbf{F}(\mathbf{U}) = \begin{bmatrix} \frac{1}{2}u_x^2 \\ \frac{1}{2}u_x u_y \end{bmatrix}, \quad \mathbf{G}(\mathbf{U}) = \begin{bmatrix} \frac{1}{2}u_x u_y \\ \frac{1}{2}u_y^2 \end{bmatrix}, \quad \mathbf{S}(x; \boldsymbol{\mu}) = \begin{bmatrix} 0.02 \exp(\mu_2 x) \\ 0 \end{bmatrix}. \quad (63)$$

The problem is defined on $\Omega = [0, 100] \times [0, 100]$ with final time $T_f = 25$ and parameters $\boldsymbol{\mu} = (\mu_1, \mu_2) \in \mathcal{D}$, where

$$\mathcal{D} = [4.25, 5.50] \times [0.015, 0.03]. \quad (64)$$

This conservative form defines the discrete residual and Jacobian used by both the PROM and HPROM pipelines. The initial/boundary conditions used in this benchmark are

$$u_x(0, y, t; \boldsymbol{\mu}) = \mu_1, \quad u_x(x, y, 0) = 1, \quad u_y(x, y, 0) = 1. \quad (65)$$

7.2 Numerical discretization

The HDM uses a cell-centered finite-volume Godunov discretization on a 250×250 mesh, hence $N = 2N_x N_y = 125,000$. Time integration uses 500 uniform steps with $\Delta t = 0.05$ (thus 501 stored states including $t = 0$). The same time-discrete residual is used by linear HPROM and by all intrusive ANN-enhanced variants (PROM-ANN and PROM-POD-AE) through LSPG/Gauss–Newton solves.

7.3 Training and evaluation protocol

Stage 1 builds $\mathbf{V}_{\text{tot}}$ from HDM snapshots at 9 structured training parameters (a 3×3 tensor grid on $\mathcal{D}$), giving $N_s = 9 \times 501 = 4509$ snapshots.

Common basis configuration. The nine trajectories provide 4500 candidate non-initial snapshot columns. For the reported basis, we retain 2250 centered snapshot columns and 2250 metric samples through the same global parameter–time-stratified rule, with uniform parameter weights and trapezoidal time weights. In Eq. (19), we set $\mathbf{P}_s = \mathbf{I}$ and choose the mass regularization by the trace-ratio rule

$$\varepsilon = \eta \frac{\text{tr}(\mathbf{C}_J)}{\text{tr}(\mathbf{C}_M)}, \quad \eta = 10^{-10}, \quad (66)$$

where $\mathbf{C}_M = \mathbf{S}_c^T \mathbf{M} \mathbf{S}_c$ and $\mathbf{C}_J$ is the sensitivity contribution to Eq. (20). This gives $\varepsilon = 6.2555 \times 10^{-10}$ for the reported data. The retained $n_{\text{tot}} = 151$ modes capture 99.9883% of the weighted snapshot energy, so the mass term regularizes the metric scale without overriding the LSPG-sensitivity contribution.

Stage 2 is performed at the same level as the subsequent evaluation: the PROM campaign uses full-residual linear LSPG trajectories, while the HPROM campaign uses linear-HPROM trajectories in the same $n_{\text{tot}} = 151$ trial space. Stage 3 then trains each learned model from the reduced-coordinate trajectories generated by the corresponding linear solver: PROM-generated trajectories in the PROM campaign and linear-HPROM trajectories in the HPROM campaign. Thus, the learning labels are not HDM snapshots subsequently projected onto the trial space. This choice keeps each learned model aligned with the coordinate dynamics of the reference solver used in its own campaign, rather than mixing full-residual or sampled-residual effects with an additional projection-label error. A conventional purely non-intrusive workflow could instead use such projected HDM labels; that alternative is not used in the present comparison.

Shared Case 2 and POD–NN master map. For the reported Case 2/POD–NN comparison, we train one full-coordinate parameter–time map per campaign,

$$\mathcal{G}_q : (\boldsymbol{\mu}, t) \mapsto \hat{\mathbf{q}}_N \in \mathbb{R}^{n_{\text{tot}}}. \quad (67)$$

This is a comparison-control decision, not a restriction of the Case 2 formulation. For a Case 2 split with n primary coordinates, the tail is the selected secondary block of this map,

$$\mathcal{M}_n(\boldsymbol{\mu}, t) = \mathbf{P}_{>n} \mathcal{G}_q(\boldsymbol{\mu}, t), \quad \mathbf{P}_{>n} \hat{\mathbf{q}}_N = \hat{\mathbf{q}} \in \mathbb{R}^{n_{\text{tot}}-n}. \quad (68)$$

Case 2 inserts $\mathcal{M}_n$ into the residual solve for the first n coordinates, whereas POD–NN-ROM uses the complete vector $\mathcal{G}_q(\boldsymbol{\mu}, t)$ directly. Therefore $n = 0$ is exactly the POD–NN-ROM prediction from the same trained map, rather than a merely similar regressor. This shared-map construction is used only to make the reported Case 2 and POD–NN results directly comparable; a Case 2 tail could otherwise be trained separately.

Primary evaluation is reported at one verification point and two off-grid points:

$$\boldsymbol{\mu}^{(v)} = (4.875, 0.0225), \quad \boldsymbol{\mu}^{(1)} = (4.56, 0.019), \quad \boldsymbol{\mu}^{(2)} = (5.19, 0.026). \quad (69)$$

For the expanded-enrichment study, we also add an extrapolatory stress-test point located outside the original training box,

$$\boldsymbol{\mu}^{(3)} = (4.00, 0.0330), \quad (70)$$

which is 20% beyond the upper-left corner of $\mathcal{D}$ in both parameter directions. Throughout the tables and figures, the superscript (v) denotes the verification point, (1) and (2) denote the two off-grid test points, and (3) denotes the extrapolatory stress-test point. Results are presented in the order $\boldsymbol{\mu}^{(v)}$, $\boldsymbol{\mu}^{(1)}$, $\boldsymbol{\mu}^{(2)}$, and $\boldsymbol{\mu}^{(3)}$ in the main baseline and enriched result tables and figures.

In addition to the baseline 3×3 training set, we consider an enriched campaign with 36 Latin-hypercube parameter points: 18 inside $\mathcal{D}$ and 18 in the 25% expanded box

$$[3.9375, 5.8125] \times [0.01125, 0.03375]. \quad (71)$$

The design has fixed seed 42. The same locations are used for the PROM and HPROM enrichment campaigns, but their coefficient targets are generated by their respective linear teachers. No additional HDM trajectory is computed; each enriched data set contains $45 \times 501 = 22,545$ coefficient snapshots. The four reported test points are excluded from both training designs. Two additional trajectories, $\boldsymbol{\mu}^{\text{val},1} = (4.5625, 0.02625)$ and $\boldsymbol{\mu}^{\text{val},2} = (5.1875, 0.01875)$, are reserved for external validation of parameter–time maps and are also distinct from the four test points.

Table 2: PROM coefficient data used for the enriched training campaign. The same $9 + 36$ parameter design is used at the HPROM level with linear-HPROM coefficient targets.

Quantity	Value
Baseline PROM trajectories	9
Interior LHS trajectories	18
Exterior LHS trajectories	18
Total training trajectories	45
Snapshots per trajectory	501
Training samples	22545
LHS seed	42
Margin fraction	0.25




Figure 1: Baseline training set and evaluation points in parameter space. The baseline consists of the 3×3 parameter grid. The axes are shared with Fig. 2, so the extrapolation point $\mu^{(3)}$ is located on the same scale as the enriched design.




Figure 2: Expanded enriched training set in parameter space. The baseline 3×3 grid is supplemented with 18 interior and 18 margin Latin-hypercube points, generated with the fixed linear teacher in the 25% expanded parameter box. The four evaluation points are excluded from both training designs. The axes match Fig. 1; the enrichment changes coefficient coverage, not the HDM snapshot basis.

7.4 PROM and HPRM Evaluation Pipelines

The PROM and HPRM comparisons are deliberately parallel but not identical. At the PROM level, the linear teacher and every intrusive learned model use the full residual. At the HPRM level, the linear teacher uses its 151-coordinate ECM rule and each learned intrusive residual uses a separate ECM rule matched to its manifold and online unknown dimension. The reported nonlinear rules use 2% global parameter–time stratified residual snapshots, direct dense SVD before ECM, and nonnegative ECM weights. The direct POD–NN–ROM and POD–DL–ROM maps do not evaluate a residual and therefore do not carry an ECM approximation error; they are reported to isolate the coefficient-learning component.

7.5 Reported Metrics

For each evaluation parameter, state accuracy is measured by the relative space–time error

$$e_u(\boldsymbol{\mu}) = 100 \frac{\|\mathbf{u}_{\text{HDM}} - \mathbf{u}_{\text{ROM}}\|_F}{\|\mathbf{u}_{\text{HDM}}\|_F}. \quad (72)$$

Coefficient diagnostics use the corresponding linear PROM or linear HPROM trajectory as reference,

$$e_q(\boldsymbol{\mu}) = 100 \frac{\|\mathbf{q}_{151}^{\text{lin}} - \hat{\mathbf{q}}_{151}\|_F}{\|\mathbf{q}_{151}^{\text{lin}}\|_F}. \quad (73)$$

We report aggregate state and coefficient errors, per-index coefficient errors, and solution overlays. The linear reference is an accuracy benchmark for the fixed trial space and residual level, not a theorem-level lower bound against the HDM.

8 Baseline PROM-Only Results

Table 3 reports the selected baseline training runs. The largest validation error occurs for the master parameter–time map because it predicts all 151 coefficients from only (μ_1, μ_2, t) . The Case 1, Case 3, and autoencoder tasks are easier because their inputs contain coefficient information.

Table 3: Baseline PROM-only Stage-3 selected networks. Errors are relative Frobenius percentages in coefficient space.

Model	Map / network	Act.	Train e_q (%)	Val. e_q (%)	Params
PROM-ANN Case 1	$\mathbf{q}_p \mapsto \mathbf{q}_s; 10 \rightarrow 256 \rightarrow 512 \rightarrow 512 \rightarrow 256 \rightarrow 141$	SILU	1.130	1.251	564,621
Master POD-NN-ROM (Case 2 tail source)	$(\mu_1, \mu_2, t) \mapsto \mathbf{q}_{151}; 3 \rightarrow 256 \rightarrow 512 \rightarrow 512 \rightarrow 256 \rightarrow 151$	ELU	2.601	4.391	565,399
PROM-ANN Case 3	$(\mathbf{q}_p, \mu_1, \mu_2, t) \mapsto \mathbf{q}_s; 13 \rightarrow 256 \rightarrow 512 \rightarrow 512 \rightarrow 256 \rightarrow 141$	SILU	0.405	0.482	565,389
PROM-POD-AE	z-score AE; $151 \rightarrow 512 \rightarrow 256 \rightarrow 128 \rightarrow 10 \rightarrow 128 \rightarrow 256 \rightarrow 512 \rightarrow 151$	GELU	0.106	0.127	486,817
POD-DL-ROM	z-score latent dynamics; $3 \rightarrow 256 \rightarrow 512 \rightarrow 512 \rightarrow 256 \rightarrow 10 \rightarrow 151$	SILU	0.810	0.935	952,747

Table 4 gives the baseline online state error. The linear PROM is not a theorem-level lower bound against the HDM, but it is the reference accuracy level associated with the fixed 151-dimensional PROM basis and teacher data. The intrusive learned PROMs are close to that reference for in-domain points. The purely non-intrusive maps degrade more strongly, especially at the extrapolatory point.

Table 4: Baseline PROM-only online errors $100 \|\mathbf{u}_{\text{HDM}} - \mathbf{u}_{\text{ROM}}\|_F / \|\mathbf{u}_{\text{HDM}}\|_F$. The mean column averages $\boldsymbol{\mu}^{(v)}, \boldsymbol{\mu}^{(1)}, \boldsymbol{\mu}^{(2)}$.

Model	$\boldsymbol{\mu}^{(v)}$	$\boldsymbol{\mu}^{(1)}$	$\boldsymbol{\mu}^{(2)}$	Mean	$\boldsymbol{\mu}^{(3)}$
Linear PROM	0.413	0.460	0.472	0.448	0.868
PROM-ANN Case 1	0.409	0.520	0.683	0.537	1.584
PROM-ANN Case 2 ($n = 10$)	1.013	1.427	1.670	1.370	1.898
PROM-ANN Case 2 ($n = 20$)	0.999	1.328	1.505	1.278	1.658
PROM-ANN Case 3	0.405	0.521	0.576	0.501	1.392
PROM-POD-AE ($n_z = 10$)	0.408	0.539	0.607	0.518	1.978
POD-NN-ROM	1.036	1.675	1.842	1.518	3.214
POD-DL-ROM ($n_z = 10$)	0.489	1.398	1.463	1.116	3.728




Figure 3: Baseline PROM solution cut-plane comparison against the HDM at the four evaluation points. Each row corresponds to one parameter, with $u_x(x, y_{\text{mid}})$ at left and $u_x(x_{\text{mid}}, y)$ at right. Fainter solid curves denote the two intermediate times and opaque solid curves the final time. All panels use the same vertical range for u_x , so the visible deviations are directly comparable across parameters and cut planes.

9 Baseline Online Coefficient Diagnostics

The online state error is the final metric, but coefficient-space diagnostics explain where the learned models lose accuracy. Table 5 reports coefficient errors from the actual online outputs. Case 1, Case 2, Case 3, and PROM-POD-AE use their PROM online trajectories, while POD-NN-ROM and POD-DL-ROM use their direct non-intrusive predictions. All coefficient errors are measured against the linear PROM trajectory with $n_{\text{tot}} = 151$. For Case 1 and Case 3, the online solver did not save $\mathbf{q}(t)$ directly, so the coefficients are recovered from the saved PROM state snapshots by

$$\mathbf{q}(t) = \arg \min_{\boldsymbol{\eta} \in \mathbb{R}^{151}} \|\mathbf{V}_{151} \boldsymbol{\eta} - (\mathbf{u}_{\text{PROM}}(t) - \mathbf{u}_{\text{ref}})\|_2. \quad (74)$$

Table 5: Baseline online coefficient error against the linear PROM with $n_{\text{tot}} = 151$. Values are relative Frobenius percentages over all 151 coefficients and all time steps.

Model	$\mu^{(v)}$	$\mu^{(1)}$	$\mu^{(2)}$	Mean	$\mu^{(3)}$
Linear PROM	0.000	0.000	0.000	0.000	0.000
PROM-ANN Case 1	0.295	0.593	1.297	0.729	2.923
PROM-ANN Case 2 ($n = 10$)	2.440	3.361	3.906	3.236	4.257
PROM-ANN Case 2 ($n = 20$)	2.401	3.123	3.512	3.012	3.662
PROM-ANN Case 3	0.164	0.856	0.788	0.603	2.396
PROM-POD-AE	0.140	0.773	0.929	0.614	4.394
POD-NN-ROM	2.499	3.924	4.249	3.557	7.444
POD-DL-ROM	0.814	3.129	3.225	2.389	8.639




Figure 4: Baseline online coefficient errors against the $n_{\text{tot}} = 151$ linear PROM trajectory. Top: absolute trajectory error $\|q_i - q_i^{\text{ref}}\|_2$. Bottom: the same error relative to $\|q_i^{\text{ref}}\|_2$, reported as a percentage. All columns in each row use the same logarithmic ordinate limits, and the vertical line marks coefficient 10. The linear PROM itself is the reference and is not plotted.

The high-index coefficients often show large relative errors because their trajectory norms are small. This does not mean that each high-index coefficient dominates the state error. It does show, however, that the map $(\mu_1, \mu_2, t) \mapsto \mathbf{q}_{151}$ is data-limited in the baseline regime: the high-index tail is low energy but still difficult to interpolate from only nine parameter trajectories.




Figure 5: Baseline PROM time-by-coefficient absolute errors $|q_i(t) - q_i^{\text{ref}}(t)|$ against the linear PROM. Columns are ordered as $\boldsymbol{\mu}^{(v)}$, $\boldsymbol{\mu}^{(1)}$, $\boldsymbol{\mu}^{(2)}$, and $\boldsymbol{\mu}^{(3)}$; rows identify the online model. The dotted horizontal line is the ten-coordinate primary/secondary split. One common colourbar is used for every panel.




Figure 6: Baseline PROM time-by-coefficient relative errors, using $100|q_i(t) - q_i^{\text{ref}}(t)|/\|q_i^{\text{ref}}\|_2$. The colour range is fixed for all PROM and HPROM heatmaps; values above 20% are shown at the upper colour limit so that lower-error structure remains visible.

9.1 Case-2 PROM Diagnostic Sweep

Case 2 connects the direct parameter-time map to the full linear PROM. We run

$$n \in \{0, 3, 5, 10, 20, 30, 50, 100, 151\}, \quad (75)$$

where $n = 0$ is exactly the POD-NN-ROM prediction from the shared master network and $n = 151$ is the linear PROM. For intermediate n , LSPG solves the first n coordinates while the same master ANN supplies the remaining coordinates.




Figure 7: Case-2 PROM state error as the number n of residual-solved primary coordinates is varied. The direct map is the $n = 0$ endpoint and the 151-coordinate linear PROM is the $n = 151$ endpoint. The black curve is the in-domain mean over $\mu^{(v)}, \mu^{(1)}, \mu^{(2)}$; the extrapolatory point $\mu^{(3)}$ remains a separate curve.

The sweep shows that solving more primary coordinates does not automatically produce monotone improvement at all points. The injected tail still affects the residual minimization, and the benefit of increasing n depends on whether the remaining predicted tail is accurate enough for that parameter point.




Figure 8: Case-2 PROM coefficient errors with respect to the linear PROM coefficient trajectories for the same n -sweep. Top: absolute time-trajectory error. Bottom: relative error. The dotted vertical line marks coefficient 10, the primary dimension used in the production Case-2 runs.

To separate the role of the injected secondary coordinates from the LSPG solve itself, we perturb the secondary coordinates around the linear PROM values for the fixed $n = 10$ split. This

diagnostic asks how much primary- coordinate and state error is induced when the secondary tail is no longer exact.

Case 2 $n=10$: sensitivity to prescribed secondary-coordinate error




Figure 9: Fixed-split sensitivity of Case-2 PROM with $n = 10$ to prescribed secondary-coordinate error. Left: state error against the HDM. Right: error in the ten solved primary coordinates against the linear PROM.

The interpretation is that Case 2 is not only a regression problem for the tail. Tail errors can change the solved primary coordinates, so the online error reflects both tail-prediction error and residual-solve feedback.

10 Effect of PROM Data Enrichment

The enriched campaign tests whether the baseline degradation is primarily an architecture issue or a data-coverage issue. The architecture is kept fixed. The basis $\mathbf{V}_{151}$, validation points, online evaluation points, and online PROM formulations are also unchanged. Only the training coefficient data are expanded from nine PROM trajectories to 45 PROM trajectories.

Table 6: Training and validation coefficient errors before and after PROM data enrichment. The validation set is kept fixed when available, so the reduction measures generalization improvement rather than a change in evaluation criterion.

Model	Base train e_q (%)	Base val. e_q (%)	Enriched train e_q (%)	Enriched val. e_q (%)	Val. reduction (%)
PROM-ANN Case 1	1.130	1.251	0.360	0.374	70.1
Master POD-NN-ROM (Case 2 tail source)	2.601	4.391	0.435	0.555	87.4
PROM-ANN Case 3	0.405	0.482	0.230	0.268	44.5
PROM-POD-AE ($n_z = 10$)	0.106	0.127	0.054	0.074	42.0
POD-DL-ROM ($n_z = 10$)	0.810	0.935	0.142	0.158	83.1

The strongest gain is obtained by the master Case 2/POD-NN map: its validation error decreases from 4.391% to 0.555%. This is the clearest evidence that the previous difficulty was largely caused by the sparsity of the nine-trajectory parameter training set, not by an intrinsic impossibility of learning the map in this benchmark. The POD-DL map also improves strongly, while the already accurate autoencoder still benefits.

Table 7: Baseline and enriched PROM state errors $100 \|\mathbf{u}_{\text{HDM}} - \mathbf{u}_{\text{ROM}}\|_F / \|\mathbf{u}_{\text{HDM}}\|_F$. Each mean averages $\mu^{(v)}, \mu^{(1)}, \mu^{(2)}$. The linear PROM is identical in both blocks because its basis and full residual formulation are fixed. The horizontal rule separates residual-evaluating PROMs from the two direct maps.

Model	Baseline (9 trajectories)					Enriched (9+36 trajectories)				
	$\mu^{(v)}$	$\mu^{(1)}$	$\mu^{(2)}$	Mean	$\mu^{(3)}$	$\mu^{(v)}$	$\mu^{(1)}$	$\mu^{(2)}$	Mean	$\mu^{(3)}$
Linear PROM	0.413	0.460	0.472	0.448	0.868	0.413	0.460	0.472	0.448	0.868
PROM-ANN Case 1	0.409	0.520	0.683	0.537	1.584	0.411	0.459	0.469	0.446	0.868
PROM-ANN Case 2 ($n = 10$)	1.013	1.427	1.670	1.370	1.898	0.435	0.477	0.496	0.469	0.872
PROM-ANN Case 3	0.405	0.521	0.576	0.501	1.392	0.412	0.458	0.470	0.447	0.869
PROM-POD-AE ($n_z = 10$)	0.408	0.539	0.607	0.518	1.978	0.412	0.459	0.471	0.447	0.871
POD-NN-ROM	1.036	1.675	1.842	1.518	3.214	0.437	0.479	0.500	0.472	0.897
POD-DL-ROM ($n_z = 10$)	0.489	1.398	1.463	1.116	3.728	0.414	0.460	0.473	0.449	0.868




Figure 10: Effect of PROM data enrichment on online state errors. Left: mean error over the verification and two off-grid in-domain points. Right: extrapolatory point.




Figure 11: Enriched PROM solution cut-plane comparison against the HDM. The same seven online models and the same shared state range as in Fig. 3 are used. Fainter solid curves are intermediate times and opaque solid curves are final-time traces; model identity is carried only by colour.

After enrichment, all learned models are close to the linear PROM state accuracy at the reported points. The most important qualitative change is that the non-intrusive maps no longer behave as outliers: the enriched POD-NN-ROM and POD-DL-ROM have state errors comparable to the intrusive learned PROMs. For the extrapolatory point, the enriched errors also collapse near the fixed linear PROM level, indicating that the exterior LHS samples provide useful information in the expanded parameter region.

Table 8: Baseline and enriched PROM coefficient errors against the corresponding fixed $n_{\text{tot}} = 151$ linear PROM coefficient trajectory. Values are relative Frobenius percentages over all coefficients and time steps. The horizontal rule separates residual-evaluating PROMs from the two direct maps.

Model	Baseline (9 trajectories)					Enriched (9+36 trajectories)				
	$\mu^{(v)}$	$\mu^{(1)}$	$\mu^{(2)}$	Mean	$\mu^{(3)}$	$\mu^{(v)}$	$\mu^{(1)}$	$\mu^{(2)}$	Mean	$\mu^{(3)}$
Linear PROM	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000
PROM-ANN Case 1	0.295	0.593	1.297	0.729	2.923	0.066	0.057	0.089	0.070	0.178
PROM-ANN Case 2 ($n = 10$)	2.440	3.361	3.906	3.236	4.257	0.377	0.357	0.433	0.389	0.583
PROM-ANN Case 3	0.164	0.856	0.788	0.603	2.396	0.039	0.038	0.055	0.044	0.117
PROM-POD-AE ($n_z = 10$)	0.140	0.773	0.929	0.614	4.394	0.041	0.041	0.059	0.047	0.151
POD-NN-ROM	2.499	3.924	4.249	3.557	7.444	0.385	0.362	0.446	0.398	0.805
POD-DL-ROM ($n_z = 10$)	0.814	3.129	3.225	2.389	8.639	0.108	0.099	0.130	0.112	0.240




Figure 12: Effect of PROM data enrichment on online coefficient errors. The coefficient metric is stricter than the state metric because it measures the entire 151-coordinate trajectory against the linear PROM teacher.




Figure 13: Enriched PROM online coefficient errors against the enriched linear-PROM reference. The plotted set is identical to Fig. 4: Case 1, Case 2 with $n = 10$, Case 3, PROM-POD-AE, POD-NN-ROM, and POD-DL-ROM. Shared ordinate ranges permit a direct baseline-enriched comparison.

The coefficient table confirms the state-error interpretation. The baseline POD–NN–ROM has an in-domain mean coefficient error of 3.557% and 7.444% at the extrapolatory point. After enrichment these decrease to 0.398% and 0.805%, respectively. Case 2 with $n = 10$ improves similarly because it uses the same master map for the injected tail.




Figure 14: Enriched PROM time-by-coefficient absolute errors against the enriched linear PROM. Rows include the four intrusive learned PROMs followed by POD–NN–ROM and POD–DL–ROM. The colour range and the coefficient-10 separator are identical to the baseline PROM heatmap in Fig. 5.




Figure 15: Enriched PROM time-by-coefficient relative errors. The shared 0–20% colour range exposes the low-error structure recovered by enrichment while clipping only larger values.




Figure 16: Case-2 per-coefficient relative error before and after PROM data enrichment for $n = 10$. The vertical line marks the separation between the ten residual-solved primary coordinates and the ANN-injected secondary tail. The linear PROM coefficient trajectory is the reference.

Figure 16 shows the mechanism relevant to the intrusive closure itself. With only nine trajectories, the ANN-injected tail has relative errors close to order one even when the global state error is moderate. The enriched training set reduces those tail errors over most coefficients, while the LSPG solve continues to recover the leading ten coordinates. This is why the Case-2 residual is far less feedback-sensitive after enrichment.

11 Baseline HPROM Results

The HPROM campaign uses the same MLSPG-sensitive basis and the same baseline nine-trajectory training set, but replaces the full residual evaluations by empirical cubature. The cubature rules are built from 2% of the available parameter-time residual snapshots using a global parameter-time stratified selection and direct dense SVD before ECM. The linear HPROM uses its own ECM rule for the 151-coordinate linear residual problem. Each learned intrusive family uses a separate rule matched to its online residual and unknown dimension. This avoids the earlier ambiguity of applying an unrelated linear rule to a nonlinear learned residual.

Table 9: Baseline HPROM cubature setup. The learned $n = 10$ rows cover Case 1, Case 2 with $n = 10$, Case 3, and PROM-POD-AE; Case 2 with $n = 20$ uses the larger $n = 20$ rule.

Family	Snapshot pct.	N_{ECM}	SVD / sampling rule
Linear HPROM	2.0	5975	direct dense SVD, global parameter-time stratified
Learned HPROM, $n = 10$	2.0	901	direct dense SVD, global parameter-time stratified
Learned HPROM, $n = 20$	2.0	1801	direct dense SVD, global parameter-time stratified

Table 10: Baseline HPROM Stage-3 selected networks. The full-coordinate master map is shared by POD-NN-ROM and Case 2; the online Case 2 solver trims this q_{151} prediction and solves the first n coordinates. Errors are relative Frobenius percentages in coefficient space.

Model	Map / network	Act.	Train e_q (%)	Val. e_q (%)	Params
HPROM-ANN Case 1	$\mathbf{q}_p \mapsto \mathbf{q}_s; 10 \rightarrow 256 \rightarrow 512 \rightarrow 512 \rightarrow 256 \rightarrow 141$	SiLU	0.992	1.101	564,621
Master POD-NN-ROM (Case 2 tail source)	$(\mu_1, \mu_2, t) \mapsto \mathbf{q}_{151}; 3 \rightarrow 256 \rightarrow 512 \rightarrow 512 \rightarrow 256 \rightarrow 151$	SiLU	0.975	1.124	565,399
HPROM-ANN Case 3	$(\mathbf{q}_p, \mu_1, \mu_2, t) \mapsto \mathbf{q}_s; 13 \rightarrow 256 \rightarrow 512 \rightarrow 512 \rightarrow 256 \rightarrow 141$	SiLU	0.491	0.577	565,389
HPROM-POD-AE	z-score AE; $151 \rightarrow 512 \rightarrow 256 \rightarrow 128 \rightarrow 10 \rightarrow 128 \rightarrow 256 \rightarrow 512 \rightarrow 151$	GELU	0.098	0.119	486,817
POD-DL-ROM	z-score latent dynamics; $3 \rightarrow 256 \rightarrow 512 \rightarrow 512 \rightarrow 256 \rightarrow 10 \rightarrow 151$	SiLU	0.624	0.744	952,747

Table 11 reports the HPROM state errors. The linear HPROM is the sampled-residual analogue of the linear PROM reference. Its in-domain mean error is 0.448%, compared with 0.868% at the extrapolatory point. Case 1, Case 3, and PROM-POD-AE remain close to this reference in the three in-domain tests, with mean errors of 0.533%, 0.494%, and 0.542%, respectively. Case 2 remains the more sensitive closure: with $n = 10$, its mean error is 1.510% and its extrapolatory error is 4.272%. Increasing the solved primary dimension to $n = 20$ reduces these values to 1.262% and 3.004%, respectively, but does not match the robustness of Case 1, Case 3, or PROM-POD-AE.

Table 11: Baseline HPROM online errors $100 \|\mathbf{u}_{\text{HDM}} - \mathbf{u}_{\text{ROM}}\|_F / \|\mathbf{u}_{\text{HDM}}\|_F$. The mean column averages $\mu^{(v)}, \mu^{(1)}, \mu^{(2)}$. The horizontal rule separates residual-evaluating HPROMs from the two direct non-intrusive maps, which do not perform an online sampled-residual solve. The corresponding ECM cardinalities are reported separately in Table 9.

Model	$\mu^{(v)}$	$\mu^{(1)}$	$\mu^{(2)}$	Mean	$\mu^{(3)}$	Mean time (s)
Linear HPROM	0.413	0.460	0.472	0.448	0.868	99.03
HPROM-ANN Case 1	0.408	0.521	0.670	0.533	1.578	18.94
HPROM-ANN Case 2 ($n = 10$)	0.505	1.904	2.121	1.510	4.272	6.03
HPROM-ANN Case 2 ($n = 20$)	0.503	1.372	1.909	1.262	3.004	9.10
HPROM-ANN Case 3	0.404	0.513	0.563	0.494	1.374	19.90
HPROM-POD-AE ($n_z = 10$)	0.417	0.548	0.662	0.542	1.846	17.32
POD-NN-ROM	0.508	1.834	2.479	1.607	7.401	0.06
POD-DL-ROM ($n_z = 10$)	0.451	1.483	1.498	1.144	4.753	0.02




Figure 17: Baseline HPROM solution cut-plane comparison against the HDM at the four evaluation points. The layout, solid-curve time-opacity convention, state range, and method colors are identical to those in Fig. 3. The linear HPROM curves are reconstructed from their saved coefficient trajectories; the nonlinear HPROM curves use the saved online state snapshots.

The overlay in Fig. 17 provides the trajectory-level counterpart to the aggregate errors. The $n = 10$ Case-2 solution is visibly less accurate at the off-grid and extrapolatory points than the linear HPROM and the other intrusive closures, although the corrected production rule no longer produces the severe departure seen in the earlier campaign. Solving twenty coordinates reduces the state error, consistent with more of the leading coefficient subspace being corrected by the online residual solve; the remaining difference is governed by the ANN-supplied secondary block.

Table 12: Baseline HPROM online coefficient error against the linear HPROM with $n_{\text{tot}} = 151$. Values are relative Frobenius percentages over all 151 coefficients and all time steps. The final two rows are the direct non-intrusive maps and are separated from the residual-evaluating HPROMs by a horizontal rule.

Model	$\mu^{(v)}$	$\mu^{(1)}$	$\mu^{(2)}$	Mean	$\mu^{(3)}$
Linear HPROM	0.000	0.000	0.000	0.000	0.000
HPROM-ANN Case 1	0.245	0.593	1.250	0.696	2.899
HPROM-ANN Case 2 ($n = 10$)	0.924	4.375	4.814	3.371	9.996
HPROM-ANN Case 2 ($n = 20$)	0.921	3.105	4.377	2.801	6.981
HPROM-ANN Case 3	0.242	0.819	0.775	0.612	2.296
HPROM-POD-AE ($n_z = 10$)	0.273	0.839	1.122	0.745	4.004
POD-NN-ROM	0.929	4.192	5.642	3.588	17.433
POD-DL-ROM ($n_z = 10$)	0.616	3.334	3.291	2.413	11.101




Figure 18: Baseline HPROM online coefficient errors against the $n_{\text{tot}} = 151$ linear HPROM trajectory. Top: absolute trajectory error. Bottom: relative trajectory error. Intrusive closures and the two direct non-intrusive maps are all included from their saved online outputs. The PROM and HPROM coefficient figures use the same fixed ordinate ranges, and the vertical line marks coefficient 10. The linear HPROM itself is the reference and is not plotted.

The HPROM coefficient diagnostics mirror the PROM-level mechanism. The high-index tail remains the difficult part of the learned map, and Case 2 is the most exposed to this difficulty because its tail is injected directly into the online residual. Increasing from $n = 10$ to $n = 20$ reduces the Case-2 coefficient error from 3.371% to 2.801% in-domain mean and from 9.996% to 6.981% at extrapolation, but it does not remove the tail limitation. Case 3 is the most accurate learned HPROM in coefficient space, with 0.612% in-domain mean error, because it combines residual-solved primary coordinates with explicit (μ_1, μ_2, t) information in the closure.




Figure 19: Baseline HPROM time-by-coefficient absolute errors against the linear HPROM. The row/column ordering, the primary-coordinate separator, and the absolute colour range are identical to those of Fig. 5.




Figure 20: Baseline HPROM time-by-coefficient relative errors against the linear HPROM. The common 0–20% colour range is exactly the one used in Fig. 6; hence colour intensities can be compared directly between the PROM and HPROM panels.

11.1 Case–2 HPROM Diagnostic Sweep

The Case–2 mechanism can be examined separately from the production comparison by varying the number of residual-solved primary coordinates. The two endpoints are the direct parameter–time map at $n = 0$ and the 151-coordinate linear HPROM at $n = 151$. For every intermediate value $n \in \{3, 5, 10, 20, 30, 50\}$, we rebuilt a matched Case–2 ECM rule from the same nine parameter trajectories using 2% global parameter–time stratified residual snapshots, seed 42, and direct dense SVD. Thus, the intermediate rows are a diagnostic of the coupled choice of primary dimension and matched cubature rule, not a continuation with one fixed sampled residual.

Table 13: Case-2 HPROM diagnostic state errors $100 \|\mathbf{u}_{\text{HDM}} - \mathbf{u}_{\text{ROM}}\|_F / \|\mathbf{u}_{\text{HDM}}\|_F$. The intermediate rows use newly built, matched 2% ECM rules; $n = 0$ and $n = 151$ are the existing direct-map and linear-HPROM references. The mean averages $\mu^{(v)}, \mu^{(1)}, \mu^{(2)}$.

Solved n	N_{ECM}	$\mu^{(v)}$	$\mu^{(1)}$	$\mu^{(2)}$	Mean	$\mu^{(3)}$
0	—	0.508	1.834	2.479	1.607	7.401
3	271	0.518	3.133	2.605	2.085	18.623
5	451	0.507	2.273	2.223	1.667	14.463
10	901	0.505	1.902	2.120	1.509	4.257
20	1801	0.504	1.372	1.909	1.261	3.005
30	2701	0.501	1.312	1.782	1.198	3.096
50	4442	0.498	1.226	1.543	1.089	2.822
151	5975	0.413	0.460	0.472	0.448	0.868




Figure 21: Case-2 HPROM state error as the number n of residual-solved primary coordinates is varied. Each intermediate point uses its matched 2% ECM rule. The dotted horizontal levels are the 151-coordinate linear-HPROM errors.

The small- n behavior is non-monotone. The direct map has an in-domain mean error of 1.607% and an extrapolatory error of 7.401%. Solving three or five coordinates is worse: the corresponding in-domain/extrapolatory pairs are 2.085%/18.623% and 1.667%/14.463%. Recovery begins with the production $n = 10$ split, which gives 1.509%/4.257%, and continues to 1.261%/3.005% at $n = 20$, 1.198%/3.096% at $n = 30$, and 1.089%/2.822% at $n = 50$. The slight extrapolatory increase from $n = 20$ to $n = 30$ confirms that the trend is not strictly monotone. All intermediate values remain above the linear-HPROM reference (0.448%/0.868%), but the $n \geq 10$ results show that the small- n degradation is not an unavoidable consequence of solving only a subset of coordinates.

Case-2 HPROM master-map sweep: coefficient errors versus linear HPROM




Figure 22: Case-2 HPROM coefficient errors for the matched-rule diagnostic sweep. Top: absolute time-trajectory error. Bottom: relative error with respect to the linear HPROM. The dotted vertical line separates the first ten coordinates. The $n = 151$ linear HPROM is the zero-error reference and is omitted.

Figure 22 shows that the residual solve chiefly reduces errors inside the solved leading block. The secondary tail remains the output of the same master map, so increasing n moves the boundary of the correction without making the learned tail exact. The resulting state trend is therefore a coupled consequence of a larger solved subspace and a different sampled residual, rather than a simple monotone interpolation between the two endpoints.

11.2 Fixed-Rule Tail-Error Sensitivity

To isolate tail feedback from the changing-rule effect, we fixed the production Case-2 $n = 10$ ECM rule ($N_e = 901$) and imposed a controlled secondary-coordinate perturbation $\mathbf{q}_s^{\text{used}} = \mathbf{q}_s^{\text{lin,HPROM}} + \alpha(\hat{\mathbf{q}}_s - \mathbf{q}_s^{\text{lin,HPROM}})$. Here $\hat{\mathbf{q}}_s$ is the master-ANN tail and α scales its actual error direction from zero to 50% of the linear-HPROM tail norm. The primary coordinates are still obtained by the same sampled LSPG solve. Stars in Fig. 23 mark the actual master-ANN tail error at each test point.

Table 14: Fixed-rule Case-2 HPROM tail diagnostic for $n = 10$; all error columns are percentages. The master-ANN tail error is measured against the linear HPROM. The zero-tail row is the constrained ten-coordinate sampled solve with the exact linear-HPROM tail; it is not the 151-coordinate linear-HPROM solution.

Point	ANN tail error	e_u at exact tail	e_u at ANN tail	e_{q_p} at ANN tail	$e_{q_{151}}$ at ANN tail
$\mu^{(v)}$	4.353	0.452	0.512	0.184	0.941
$\mu^{(1)}$	14.695	0.564	4.516	10.164	10.420
$\mu^{(2)}$	21.740	0.475	2.895	4.865	6.538
$\mu^{(3)}$	38.799	1.077	13.098	30.445	30.777




Figure 23: Fixed-rule sensitivity of Case-2 HPROM with $n = 10$ to secondary-coordinate error. Left: state error against the HDM. Right: error in the ten solved primary coordinates against the linear HPROM. Stars identify the actual master-ANN tail errors.

The fixed-rule result establishes direct residual feedback. Even an exact linear-HPROM tail leaves a nonzero zero-tail baseline because a constrained ten-coordinate sampled solve is not the full 151-coordinate linear-HPROM problem. Additional tail error then changes the solved primary coordinates: at off-grid point $\mu^{(1)}$, a 14.695% master-ANN tail error produces a 10.164% primary-coordinate error and a 4.516% state error; at the extrapolatory point, a 38.799% tail error produces 30.445% and 13.098%, respectively. Hence Case 2 is not a passive reconstruction map: its learned tail enters the online residual and can shift the coordinates that are nominally solved intrusively.

12 Effect of HPROM Coefficient-Data Enrichment

The HPROM enrichment campaign uses the same 36 Latin-hypercube parameter locations as the PROM study: 18 interior points and 18 points in the 25% expanded exterior rectangle, with seed 42. The fixed basis, the 2% parameter-time sampling ratio, direct dense SVD, and the nine baseline parameter locations used to construct each nonlinear ECM rule are retained. The nonlinear residual maps do change with the retrained neural networks, so a rule is rebuilt for each enriched learned model, but no new HDM trajectory or residual-training parameter location is introduced. Consequently, this experiment isolates the effect of enlarging the linear-HPROM coefficient teacher data from 9 to 45 trajectories.

Table 15: Reported coefficient-space training and validation errors for the HPROM learned maps before and after coefficient-data enrichment. Values are relative Frobenius percentages. The baseline records use their original internal held-out snapshot split; the enriched master-map record uses the two separate validation trajectories, so the validation columns are useful diagnostics but are not a strictly paired model-selection protocol.

Model	Base train e_q (%)	Base val. e_q (%)	Enriched train e_q (%)	Enriched val. e_q (%)	Val. reduction (%)
HPROM-ANN Case 1	0.992	1.101	0.358	0.373	66.1
Master POD-NN-ROM (Case 2 tail source)	0.975	1.124	0.404	0.489	56.5
HPROM-ANN Case 3	0.491	0.577	0.202	0.234	59.4
HPROM-POD-AE ($n_z = 10$)	0.098	0.119	0.049	0.066	45.0
POD-DL-ROM ($n_z = 10$)	0.624	0.744	0.130	0.142	80.9

The training diagnostics improve for every intrinsic closure. For example, the reported Case 1 validation error decreases from 1.101% to 0.373%, Case 3 decreases from 0.577% to 0.234%,

and the master parameter–time map decreases from 1.124% to 0.489%. The online comparisons below are the decisive test, because all four points are kept out of both training campaigns.

Table 16: Baseline and enriched state errors for the learned models trained on linear-HPROM coefficient data. $100 \|\mathbf{u}_{\text{HDM}} - \mathbf{u}_{\text{ROM}}\|_F / \|\mathbf{u}_{\text{HDM}}\|_F$. Each mean averages $\boldsymbol{\mu}^{(v)}, \boldsymbol{\mu}^{(1)}, \boldsymbol{\mu}^{(2)}$. The linear HPROM is identical in both columns because its basis and ECM rule are fixed. The horizontal rule separates residual-evaluating HPROMs from the two non-intrusive direct maps.

Model	Baseline (9 trajectories)					Enriched (9+36 trajectories)				
	$\mu^{(v)}$	$\mu^{(1)}$	$\mu^{(2)}$	Mean	$\mu^{(3)}$	$\mu^{(v)}$	$\mu^{(1)}$	$\mu^{(2)}$	Mean	$\mu^{(3)}$
Linear HPROM	0.413	0.460	0.472	0.448	0.868	0.413	0.460	0.472	0.448	0.868
HPROM-ANN Case 1	0.408	0.521	0.670	0.533	1.578	0.413	0.459	0.471	0.448	0.870
HPROM-ANN Case 2 ($n = 10$)	0.505	1.904	2.121	1.510	4.272	0.422	0.465	0.485	0.457	0.912
HPROM-ANN Case 3	0.404	0.513	0.563	0.494	1.374	0.410	0.454	0.470	0.445	0.875
HPROM-POD-AE ($n_z = 10$)	0.417	0.548	0.662	0.542	1.846	0.435	0.478	0.488	0.467	0.874
POD-NN-ROM	0.508	1.834	2.479	1.607	7.401	0.423	0.466	0.487	0.458	0.926
POD-DL-ROM ($n_z = 10$)	0.451	1.483	1.498	1.144	4.753	0.414	0.460	0.473	0.449	0.869

Table 17: HPROM online wall time and speed-up relative to the fixed HDM reference $T_{\text{HDM}} = 737.44\text{s}$, the mean wall time of the nine HDM training trajectories. Intrusive times are means over the four evaluation trajectories and exclude offline ECM construction. The linear HPROM is the same fixed basis and ECM rule in both campaigns; its baseline mean is therefore reproduced in the enriched column. The horizontal rule separates sampled-residual HPROMs from the two direct non-intrusive maps. Direct-map times are means from ten repeated loaded-network forward predictions at each evaluation point, excluding checkpoint loading, state reconstruction, HDM evaluation, and file I/O.

Model	Baseline (9 trajectories)		Enriched (9+36 trajectories)	
	Mean time (s)	T_{HDM}/T	Mean time (s)	T_{HDM}/T
Linear HPROM	99.03	7.4	99.03	7.4
HPROM-ANN Case 1	18.94	38.9	19.44	37.9
HPROM-ANN Case 2 ($n = 10$)	6.03	122.2	5.99	123.0
HPROM-ANN Case 2 ($n = 20$)	9.10	81.0	8.57	86.1
HPROM-ANN Case 3	19.90	37.1	19.54	37.7
HPROM-POD-AE ($n_z = 10$)	17.32	42.6	18.06	40.8
POD-NN-ROM	0.0079	92,966.3	0.0076	97,669.1
POD-DL-ROM ($n_z = 10$)	0.0103	71,823.6	0.0100	73,414.5

The sampled-residual closures provide a substantial online reduction relative to the HDM. For Case-2 with $n = 10$, the mean time is 6.03 s in the baseline campaign and 5.99 s after enrichment, corresponding to $122.2\times$ and $123.0\times$ speed-ups relative to the fixed HDM reference. The direct POD-NN-ROM and POD-DL-ROM maps do not solve a residual minimization problem; their entries in Table 17 provide computational context rather than an ECM comparison.




Figure 24: Effect of HPROM coefficient-data enrichment on state error. The black level is the fixed linear HPROM reference; logarithmic axes make both the baseline Case-2 degradation and the enriched recovery visible.




Figure 25: Enriched HPROM solution cut-plane comparison against the HDM. The same seven online models and the same shared state range as in Fig. 17 are used. Fainter solid curves are intermediate times and opaque solid curves are final-time traces; model identity is carried only by colour.

Coefficient-data enrichment largely eliminates the state-level Case-2 gap. At $n = 10$, its in-domain mean state error falls from 1.510% to 0.457%, and the extrapolatory error falls from 4.272% to 0.912%. The latter is close to the 0.868% linear-HPROM reference. Case 1 and Case 3 likewise return to the fixed linear HPROM level, with enriched in-domain means of 0.448% and 0.445%, respectively. PROM-POD-AE remains accurate, at 0.467% in-domain mean and 0.874% at extrapolation. The production paired comparison keeps $n = 10$ as the Case-2 closure; the $n = 20$ run is retained as a primary-dimension diagnostic rather than treated as a separate learned model family.

Table 18: Baseline and enriched coefficient errors against the corresponding fixed $n_{\text{tot}} = 151$ linear HPROM coefficient trajectory. Values are relative Frobenius percentages over all coefficients and all time steps. The horizontal rule separates residual-evaluating HPROMs from the two non-intrusive direct maps.

Model	Baseline (9 trajectories)					Enriched (9+36 trajectories)				
	$\mu^{(v)}$	$\mu^{(1)}$	$\mu^{(2)}$	Mean	$\mu^{(3)}$	$\mu^{(v)}$	$\mu^{(1)}$	$\mu^{(2)}$	Mean	$\mu^{(3)}$
Linear HPROM	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000
HPROM-ANN Case 1	0.245	0.593	1.250	0.696	2.899	0.072	0.067	0.090	0.076	0.167
HPROM-ANN Case 2 ($n = 10$)	0.924	4.375	4.814	3.371	9.996	0.342	0.305	0.398	0.348	0.681
HPROM-ANN Case 3	0.242	0.819	0.775	0.612	2.296	0.055	0.064	0.059	0.059	0.121
HPROM-POD-AE ($n_z = 10$)	0.273	0.839	1.122	0.745	4.004	0.273	0.288	0.254	0.272	0.316
POD-NN-ROM	0.929	4.192	5.642	3.588	17.433	0.341	0.300	0.396	0.345	0.752
POD-DL-ROM ($n_z = 10$)	0.616	3.334	3.291	2.413	11.101	0.096	0.093	0.116	0.102	0.243




Figure 26: Effect of HPROM coefficient-data enrichment on global coefficient error. Enrichment reduces the Case-2 tail error by more than one order of magnitude at the extrapolatory point, while retaining sub-percent global coefficient errors for the other intrusive closures.




Figure 27: Enriched HPROM online coefficient errors against the enriched linear-HPROM reference. The plotted set is identical to Fig. 18: Case 1, Case 2 with $n = 10$, Case 3, PROM-POD-AE, POD-NN-ROM, and POD-DL-ROM. Shared ordinate ranges permit a direct baseline-enriched comparison.

Coefficient space confirms that the improvement is not merely cancellation in the state norm. The Case-2 $n = 10$ in-domain mean coefficient error decreases from 3.371% to 0.348%, and its extrapolatory coefficient error decreases from 9.996% to 0.681%. The linear HPROM reference is unchanged; the observed improvement therefore comes from coefficient coverage rather than a different linear basis or a more accurate linear cubature rule. The $n = 20$ variant remains in the baseline comparison and the dedicated Case-2 diagnostic; the paired enrichment figures retain the production $n = 10$ closure for a one-to-one comparison with the PROM study.

Direct maps trained on linear-HPROM data. POD-NN-ROM and POD-DL-ROM do not evaluate a residual and therefore are not HPROM solvers; they have no ECM approximation error. Their baseline and enriched values are included in the final two rows of Tables 16 and 18, below a horizontal rule, rather than being reported in a separate table. This keeps the direct maps beside the intrusive models trained on the same coefficient data while preserving the distinction between a sampled-residual HPROM and a pure regression map. The HPROM Case 2 closure uses the very same POD-NN master map for its injected secondary block, so its comparison with the direct map isolates the consequence of residual correction rather than a change of regressor.

The baseline direct maps expose the same coverage limitation as Case 2. The POD-NN-ROM in-domain mean state error falls from 1.607% to 0.458% after enrichment, and its extrapolatory error falls from 7.401% to 0.926%. The corresponding POD-DL-ROM values fall from 1.144% to 0.449% in-domain and from 4.753% to 0.869% at extrapolation. Thus, the improvement is due entirely to coefficient-data coverage, not to a sampled residual solve. Their state and coefficient reductions are reported in the same paired tables and all-model bar figures as the residual-evaluating methods, ruling out an interpretation based only on cancellation in the state norm.




Figure 28: Enriched HPROM time-by-coefficient absolute errors against the enriched linear HPROM. Rows include the four intrusive learned HPROMs followed by POD–NN–ROM and POD–DL–ROM. The colour range and coefficient-10 separator are identical to the baseline HPROM heatmap in Fig. 19.




Figure 29: Enriched HPROM time-by-coefficient relative errors. The shared 0–20% colour range exposes the low-error structure recovered by enrichment while clipping only larger values.




Figure 30: Case-2 per-coefficient HPROM relative error before and after coefficient-data enrichment for $n = 10$. The dotted vertical line marks the separation between the ten residual-solved primary coordinates and the ANN-injected secondary tail. The linear HPROM coefficient trajectory is the reference.

Figure 30 gives the mechanism behind the aggregate results. In the baseline, the injected tail is inaccurate at most secondary coordinates, especially away from the training grid. Enrichment substantially lowers those tail errors. The sampled residual solve then needs only to correct the leading coordinates, which is why even Case 2 with $n = 10$ returns to the linear-HPROM state-accuracy scale.

13 Conclusions

The PROM and HPROM studies support the following conclusions.

1. The fixed MLSPG-sensitive linear PROM provides a reproducible coefficient teacher and state-accuracy reference for this benchmark.
2. With only nine training trajectories, direct parameter–time prediction of all 151 coefficients is the limiting task. This affects both the POD–NN–ROM and the Case 2 injected tail.
3. Solving more primary coordinates in Case 2 can help, but the matched-rule HPROM sweep is non-monotone at small n : $n = 3$ and $n = 5$ worsen the extrapolatory error to 18.623% and 14.463%, whereas $n = 50$ reduces it to 2.822% without reaching the 0.868% linear-HPROM reference. With a fixed $n = 10$ rule, imposed tail error changes the primary coordinates recovered by the residual solve; tail accuracy is therefore essential.
4. PROM data enrichment changes the picture substantially. Keeping the architecture and basis fixed, the 9+36 PROM trajectory set reduces validation and online coefficient errors and brings the learned state errors close to the linear PROM level.
5. The baseline HPROM with 2% parameter–time stratified snapshots preserves the main PROM conclusion for Case 1, Case 3, and PROM–POD–AE: these learned intrusive models remain close to the linear HPROM reference at in-domain points.
6. Case 2 is the residual-sampling stress test. With nine coefficient trajectories, $n = 10$ is highly sensitive to the injected secondary tail. The fixed-rule diagnostic demonstrates that this sensitivity is residual feedback, not solely reconstruction error, while the matched-rule sweep shows improvement only once a sufficiently large primary block is solved.
7. HPROM coefficient-data enrichment transfers the PROM-level gain without changing the linear HPROM reference or adding HDM training snapshots. The 9 + 36 campaign reduces Case–2 $n = 10$ from 1.510% to 0.457% in-domain and from 4.272% to 0.912% at extrapolation, placing it on the same practical accuracy scale as the linear HPROM.
8. The same enrichment effect is present without an online residual solve. Both direct maps improve sharply, particularly at extrapolation, confirming that coefficient-data coverage is a primary limitation. These maps do not quantify ECM error, however, so they complement rather than replace the intrusive HPROM comparisons.

A Matched Case–2 PROM–HPROM Diagnostics

The main results deliberately separate the full-residual PROM from the ECM-hyperreduced HPROM. This appendix makes that separation explicit for the Case–2 parameter–time closure with $n = 10$ solved primary coordinates. The comparison is diagnostic rather than a claim that hyperreduction produces a single uniform error penalty. Both state errors are measured against the same HDM trajectory. In contrast, each coefficient error is measured against the full linear reference associated with the corresponding solver: the linear PROM for the solid curves and

the linear HPROM for the dashed curves. This is the consistent reference for the prescribed-tail experiment, whose perturbation is itself defined relative to the respective linear tail.

Figure 31 compares the response to the same imposed secondary-coordinate discrepancy. The actual master-map tail levels are marked separately because they are not the same for PROM and HPROM. At the verification parameter, the zero-tail states are close, but the HPROM is substantially more sensitive once the comparison is made at a common disturbance level. For example, at a 50% imposed tail error, the state error is 12.35% for HPROM versus 5.31% for PROM, while the error in the ten solved coordinates is 27.87% versus 7.82%, respectively. The same qualitative amplification is visible at the two off-grid points and at extrapolation. Thus, this experiment isolates an additional residual-sampling sensitivity of the fixed ECM rule; it does not attribute the entire Case-2 error to the neural map.

Case-2, $n = 10$: matched PROM-HPROM secondary-tail sensitivity




Figure 31: Matched Case-2 secondary-tail diagnostic with $n = 10$. Rows are, from top to bottom, the verification point, the two off-grid points, and the extrapolatory point. Solid curves are the full-residual PROM and dashed curves are the ECM-HPROM. Both use the same trained parameter-time map; the star and cross mark the actual PROM and HPRM master-map tail errors, respectively. State error is measured against the HDM in both cases. Each primary-coordinate error is measured against the method-matched full linear reference trajectory.

Figure 32 gives the requested coefficient-resolved comparison of the two matched-rule sweeps. Each colour identifies the same solved dimension n , with PROM shown by a solid line and HPRM by a dashed line. Only the common values $n \in \{0, 3, 5, 10, 20, 30, 50\}$ are shown because the HPRM diagnostic was run only through $n = 50$; the $n = 151$ linear references are zero-error endpoints and are omitted. The verification panel does not by itself support a generic HPRM-failure claim. The systematic discrepancy instead appears at the off-grid and extrapolatory points, especially in the injected-tail range $i > 10$, where the dashed HPRM

curves remain well above the solid PROM curves for the same n . This is the coefficient-level counterpart of the fixed-tail sensitivity observed in Figure 31.




Figure 32: Matched Case-2 coefficient-error sweeps. Columns are ordered as $\mu^{(v)}$, $\mu^{(1)}$, $\mu^{(2)}$, and $\mu^{(3)}$. Top: absolute time-trajectory coefficient error. Bottom: relative coefficient error. Each colour identifies the solved primary dimension n ; solid curves are PROM and dashed curves are HPROM. The vertical line marks coefficient 10. Each solver is compared with its own full linear reference trajectory, so the figure exposes the effect of hyperreduction on the matched Case-2 coefficient sweep without conflating the two distinct linear references.

B Burgers Basis-Selection Diagnostics

The M -regularized LSPG-sensitive basis was selected through two diagnostics. The first diagnostic is the low/high transfer operator $\mathbf{T}_{LH}$ introduced in (23). Table 19 reports the singular-value statistics of this operator over the three Burgers evaluation trajectories. The LSPG-sensitive basis reduces the mean value of $\sigma_{\max}(\mathbf{T}_{LH})$ by approximately one half relative to the Euclidean POD basis, and also reduces the 95th percentile and maximum values. This supports the use of the modified metric as a robustness device for intrusive nonlinear reduced models: if the nonlinear closure, decoder, or surrogate produces errors in the unresolved part of the representation, the resulting feedback into the resolved coordinates through the LSPG normal equations is less amplifying.

Table 19: Burgers low/high transfer diagnostic for basis selection. The table reports statistics of $\sigma_{\max}(\mathbf{T}_{LH})$ and of the corresponding M -weighted state gain g_M over the original three evaluation trajectories.

Basis	mean $\sigma_{\max}(T_{LH})$	p95 $\sigma_{\max}(T_{LH})$	max $\sigma_{\max}(T_{LH})$	mean g_M	p95 g_M	max g_M
Euclidean POD	3.078e-02	3.668e-02	4.088e-02	3.078e-02	3.668e-02	4.088e-02
LSPG-sensitive POD	1.548e-02	2.244e-02	2.433e-02	1.599e-02	2.339e-02	2.559e-02

As a complementary check, we perturb only the high-coordinate block of an oracle reconstruction and measure the induced deviation from the linear reduced trajectory. This test is implemented with the Case 2 partition because it isolates the most direct contamination

mechanism: the high block is injected independently of the online low-coordinate solve. Table 20 reports averages over the original verification point, two off-grid points, and five random seeds for the perturbed cases. The M -regularized LSPG-sensitive basis slightly increases the unperturbed linear error in this preliminary test, but reduces the high-to-low contamination measured against the linear reference, especially in the resolved coordinates. The intended role of the metric is broader than this specific test: it provides a protective basis choice for nonlinear intrusive models in which learned high coordinates or decoded manifold corrections may contain approximation error. This diagnostic is not used as a standalone performance result; it supports the basis choice used in the main Burgers campaign.

Table 20: Oracle high-mode perturbation diagnostic for basis selection. Errors are averaged over $\boldsymbol{\mu}^{(v)}$, $\boldsymbol{\mu}^{(1)}$, $\boldsymbol{\mu}^{(2)}$; perturbed rows also average over five random seeds.

Basis metric	Unperturbed ε_u vs HDM (%)	1% pert., ε_u vs linear (%)	1% pert., $\varepsilon_q^{\text{low}}$ vs linear (%)	2% pert., ε_u vs linear (%)	2% pert., $\varepsilon_q^{\text{low}}$ vs linear (%)
Euclidean POD	0.441	0.091	0.014	0.183	0.045
M -regularized LSPG-sensitive POD	0.448	0.088	0.012	0.177	0.035

References

- [1] Athanasios C Antoulas. *Approximation of large-scale dynamical systems*. SIAM, 2005.
- [2] Alfio Quarteroni, Gianluigi Rozza, et al. *Reduced order methods for modeling and computational reduction*, volume 9. Springer, 2014.
- [3] Joshua Barnett, Charbel Farhat, and Yvon Maday. Neural-network-augmented projection-based model order reduction for mitigating the Kolmogorov barrier to reducibility. *Journal of Computational Physics*, 492:112420, 2023.
- [4] Albert Cohen, Charbel Farhat, Yvon Maday, and Agustin Somacal. Nonlinear compressive reduced basis approximation for PDEs. *Comptes Rendus. Mécanique*, 351(S1):357–374, 2023.
- [5] Joubine Aghili, Hassan Ballout, Yvon Maday, and Christophe Prud’Homme. Nonlinear compressive reduced basis approximation: when taylor meets kolmogorov. *arXiv preprint arXiv:2601.13712*, 2026.
- [6] Francesco Romor, Giovanni Stabile, and Gianluigi Rozza. Explicable hyper-reduced order models on nonlinearly approximated solution manifolds of compressible and incompressible navier-stokes equations. *Journal of Computational Physics*, 524:113729, 2025.
- [7] Stefania Fresca and Andrea Manzoni. Pod-dl-rom: Enhancing deep learning-based reduced order models for nonlinear parametrized pdes by proper orthogonal decomposition. *Computer Methods in Applied Mechanics and Engineering*, 388:114181, 2022.
- [8] Joaquin Alberto Hernandez, Manuel Alejandro Caicedo, and Alex Ferrer. Dimensional hyper-reduction of nonlinear finite element models via empirical cubature. *Computer Methods in Applied Mechanics and Engineering*, 313:687–722, 2017.
- [9] JA Hernández. A multiscale method for periodic structures using domain decomposition and ecm-hyperreduction. *Computer Methods in Applied Mechanics and Engineering*, 368:113192, 2020.
- [10] S Ares De Parga, Jose Raul Bravo, JA Hernández, Rubén Zorrilla, and Riccardo Rossi. Hyper-reduction for petrov–galerkin reduced order models. *Computer Methods in Applied Mechanics and Engineering*, 416:116298, 2023.

- [11] Sebastian Ares De Parga, Radek Tezaur, Carlos G. Hernández, and Charbel Farhat. Non-linear projection-based model order reduction with machine learning regression for closure error modeling in the latent space. *Computer Methods in Applied Mechanics and Engineering*, 448:118443, 2026.
- [12] Charbel Farhat, Todd Chapman, and Philip Avery. Structure-preserving, stability, and accuracy properties of the energy-conserving sampling and weighting method for the hyper reduction of nonlinear finite element dynamic models. *International Journal for Numerical Methods in Engineering*, 102(5):1077–1110, 2015.
- [13] Matthew R. Chmiel, Joshua L. Barnett, and Charbel Farhat. Assessment of projection-based model order reduction for a benchmark hypersonic flow problem. In *AIAA SciTech 2024 Forum*, page 0250, 2024.
- [14] Matthew R Chmiel, Joshua Barnett, and Charbel Farhat. Unified LSPG model reduction framework and assessment for hypersonic computational fluid dynamics. *AIAA Journal*, 63(1):72–90, 2025.